

How to interpret the evidence

Evidence	What it shows	Interpretation limit	Audit response
Clean baseline	Raw attention shows little advantage over matched-random subsets	This does not mean that every attention map is wrong	Test decision relevance
Stale exposure	More attention is attached to old vehicle data	This does not show that AoI causes lower faithfulness	Show freshness separately
Within-policy association	Greater stale exposure tends to accompany lower DEF	Association is not causation or dispatch-only evidence	Check both action groups
Direct paired change	The effect varies across trained checkpoints	There is no single response to telemetry degradation	Audit every checkpoint

Current decision: WITHHOLD raw attention as an explanation