

How to interpret the evidence

Evidence	What it shows	Interpretation limit	Audit response
Clean baseline	Raw attention shows no reliable matched-random advantage	A plausible attention map is not evidence of faithfulness	Test decision relevance
Stale exposure	Longer outages increase the measured stale-data exposure	This does not show that AoI reduces faithfulness	Show freshness separately
Within-policy association	DEF declines for only a minority of checkpoints	The proposed relationship is checkpoint dependent	Report the mixed evidence
Direct paired change	Attention reallocation varies across trained checkpoints	A single degradation effect cannot be claimed	Audit every checkpoint

Current decision: WITHHOLD raw attention as an explanation