

When Explanations Outlive Their Data: Faithfulness Decoupling in Graph-Attention MARL Fleet Dispatch under Telemetry Degradation

A Master's Dissertation

submitted in partial fulfilment of the requirements for the award of the degree of

MSc in Artificial Intelligence

By

Hongwei Lin
P2982757

Supervisor: Farhan S. Ujager



School of Computing
De Montfort University, Dubai

September 2026

DECLARATION

I declare that this dissertation is my own original work carried out under the supervision of Farhan S. Ujager and that all sources and references have been acknowledged appropriately. This dissertation has not been submitted previously for academic credit at De Montfort University, Dubai, or any other institution.

Signature: _____

Name: Hongwei Lin

Date: _____

ACKNOWLEDGMENT

I would like to express my sincere gratitude to my supervisor, Dr Farhan S. Ujager, for his valuable support and guidance throughout this research. His advice, encouragement and careful feedback helped me refine the experiment, interpret the findings responsibly and complete this dissertation.

Student signature: _____

Student name: Hongwei Lin

LIST OF FIGURES

Figure 1.1	Stale data and the explanation-release decision.....	3
Figure 3.1	Local observation graph and request-to-action mapping.....	11
Figure 3.2	Graph-attention policy and audit channel.....	13
Figure 3.3	Training, selection and held-out evaluation.....	15
Figure 3.4	Tunnel-triggered observation-layer telemetry degradation.....	16
Figure 3.5	Construct-validity controls for action-linked request nodes.....	18
Figure 4.1	GAT and GAT-Outage training and checkpoint-selection diagnostics.....	25
Figure 4.2	Clean-telemetry pickups by training seed.....	26
Figure 4.3	Faithfulness perturbation controls by checkpoint.....	28
Figure 4.4	Attention aggregation sensitivity.....	29
Figure 4.5	Attention query-row sensitivity.....	30
Figure 4.6	Top-k overlap and DEF resolution.....	31
Figure 4.7	Outage-duration sweep by policy and training seed.....	32
Figure 4.8	Paired attention and faithfulness shift by training seed.....	33
Figure 4.9	Faithfulness analysis by action type.....	34
Figure 4.10	Tunnel-triggered and random telemetry-loss sensitivity.....	36
Figure 4.11	WAMSN-DEF relationship and H4 result.....	37
Figure 5.1	Cross-seed evidence matrix.....	40
Figure 5.2	Freshness-aware explanation audit workflow.....	44

LIST OF TABLES

Table 2.1	Recent graph-based dispatch literature compared with this study.....	9
Table 3.1	Observation graph node types and roles.....	11
Table 3.2	Policy conditions used in the experiment.....	13
Table 3.3	Validation-selected checkpoint indices (zero-based).....	13
Table 3.4	Training and optimization settings.....	14
Table 3.5	Checkpoint-acceptance criteria.....	14
Table 3.6	Explanation measures and roles.....	17
Table 3.7	Held-out evaluation structure and episode counts.....	19
Table 3.8	Telemetry conditions used in held-out evaluation.....	19
Table 3.9	Hypothesis map and analysis roles.....	21
Table 4.1	Policy capability on held-out demand.....	26
Table 4.2	Construct-validity audit of random controls.....	27
Table 4.3	Clean-telemetry faithfulness controls by training seed.....	27
Table 4.4	Valid node-count distribution in scored graphs.....	30
Table 4.5	Exposure-conditioned paired attention and DEF shifts.....	32
Table 4.6	Action composition and dispatch-only DEF.....	33
Table 4.7	Tunnel-triggered and random-loss sensitivity results.....	36
Table 4.8	Hypothesis outcomes across training seeds.....	37
Table 4.9	Freshness-aware explanation audit decisions.....	38
Table 5.1	Freshness-aware explanation audit framework.....	43
Table 5.2	Audit framework inputs and outputs.....	44

LIST OF ABBREVIATIONS

Abbreviation	Full form
AoI	Age of Information
BMG-Q	Localized Bipartite Match Graph Attention Q-learning
CI	Confidence Interval
CTDE	Centralized Training with Decentralized Execution
DEF	Decision-level Explanation Faithfulness
GAT	Graph Attention Network
GNN	Graph Neural Network
GxI	Gradient x Input
LOO	Leave-One-Out
MADDPG	Multi-Agent Deep Deterministic Policy Gradient
MAPPO	Multi-Agent Proximal Policy Optimization
MARL	Multi-Agent Reinforcement Learning
MLP	Multilayer Perceptron
PPO	Proximal Policy Optimization
RL	Reinforcement Learning
SUMO	Simulation of Urban Mobility
WAMSN	Weighted Attention Mass on Stale Nodes
XRL	Explainable Reinforcement Learning

ABSTRACT

Graph-attention multi-agent reinforcement learning (MARL) can expose attention weights as explanations of fleet-dispatch decisions. However, an attention map may remain visible after its supporting vehicle telemetry becomes stale. This dissertation asks whether raw graph-attention weights can be trusted as explanations under clean and degraded telemetry.

The experiment simulates 20 taxis in Simulation of Urban Mobility (SUMO). Tunnel entry triggers an observation-layer outage that freezes an affected taxi's last valid position and speed for 10, 20, 30, or 60 seconds while SUMO retains the current traffic state. This produces paired clean and degraded observations for the same decision. Graph-attention network (GAT) policies trained on clean data or with 30-second outages are evaluated across three independent training seeds. Decision-level explanation faithfulness (DEF) compares attention-ranked nodes with type-matched, action-protected random controls. Weighted attention mass on stale nodes (WAMSN) separately measures attention assigned to stale vehicle information. Further checks cover leave-one-out rankings, attention extraction, action type, and random telemetry-loss triggers.

The results do not validate raw attention as a reliable explanation. Clean corrected DEF remains near zero and varies across checkpoints. Longer outages increase stale-data exposure, but paired attention and DEF shifts reverse direction for one training seed. The main associations are dominated by no-op decisions and do not reproduce for dispatch actions. The seed-dependent pattern also appears under random triggering and in both training regimes. Therefore, the tested attention weights do not provide consistent evidence of freshness or decision relevance.

The dissertation proposes a freshness-aware explanation audit framework that reports freshness and action composition separately, applies action-aware faithfulness tests, checks the attention extraction rule, and requires consistent evidence across checkpoints. Attention remains an internal diagnostic when these checks fail. The framework governs explanation release; it does not repair the model. Its demonstrative application returns WITHHOLD for both model families and all six checkpoints.

Keywords: explainable reinforcement learning; graph attention; multi-agent reinforcement learning; fleet dispatch; telemetry degradation; faithfulness.

TABLE OF CONTENTS

Chapter 1: Introduction.....	1
1.1 Context.....	1
1.2 Problem statement.....	1
1.3 Research questions.....	1
1.4 Objectives.....	2
1.5 Main findings.....	2
1.6 Contributions.....	3
1.7 Dissertation structure.....	4
Chapter 2: Literature Review.....	5
2.1 Reinforcement learning for fleet dispatch.....	5
2.2 Explainability in reinforcement learning.....	5
2.3 Is attention an explanation?.....	6
2.4 Explaining graph neural networks.....	7
2.5 Evaluating faithfulness: methods and pitfalls.....	7
2.6 Age of Information and telemetry degradation.....	8
2.7 From explanation evaluation to explanation assurance.....	8
2.8 Research gap and positioning.....	8
Chapter 3: Methodology.....	10
3.1 Study design.....	10
3.2 SUMO environment and data.....	10
3.3 Observation graph and action space.....	10
3.4 Policy models and training.....	11
3.5 Telemetry degradation.....	15
3.6 Explanation measures.....	16
3.6.1 Decision-level explanation faithfulness.....	16
3.6.2 Stale-node attention.....	17
3.7 Construct-validity audit.....	18
3.8 Evaluation design.....	19
3.8.1 Primary evaluation.....	19
3.8.2 Supporting analyses.....	21
3.9 Hypotheses and statistical analysis.....	21
3.9.1 Statistical procedure.....	22
3.9.2 Interpretation limits.....	22
3.10 Audit decision implementation.....	22
3.11 Reproducibility.....	23
3.12 Ethics and data governance.....	24
Chapter 4: Results.....	25
4.1 Policy capability and training stability.....	25
4.2 Construct-validity audit.....	27
4.3 Evaluator sensitivity and ranking controls.....	27
4.4 Small-graph resolution.....	30
4.5 Telemetry manipulation and stale exposure.....	31
4.6 Exposure-conditioned paired audit.....	32

4.7 Analysis by action type.....	33
4.8 Random-trigger sensitivity analysis.....	34
4.9 Faithfulness hypotheses.....	37
4.10 Result summary.....	38
4.11 Demonstrative framework decision.....	38
Chapter 5: Discussion.....	40
5.1 Answer to the central problem.....	40
5.2 What the stale-exposure result means.....	41
5.3 Dependence on checkpoint and analysis choice.....	41
5.4 Role of the construct-validity audit.....	42
5.5 Why outage training did not solve the problem.....	42
5.6 Proposed audit framework.....	43
5.6.1 Purpose and scope.....	43
5.6.2 Inputs, outputs and use.....	44
5.6.3 Application to this study.....	45
5.7 Limitations.....	45
5.8 Future work.....	46
Chapter 6: Conclusion.....	47

Chapter 1: Introduction

1.1 Context

Fleet dispatch is relational. A taxi's useful action depends on nearby vehicles, open passenger requests, road conditions, and competition for the same demand. Graph neural networks (GNNs) represent these relationships directly, and graph attention networks (GATs) add learned weights over neighboring nodes (Scarselli et al., 2009; Velickovic et al., 2018). Because the weights are easy to display, they are often read as an explanation of which vehicles or requests influenced a decision.

That reading is convenient, but it is not automatically valid. Prior work has shown that an attention weight can be weakly related to the evidence that actually changes a model's output (Jain and Wallace, 2019; Serrano and Smith, 2019; Wiegrefe and Pinter, 2019; Liu et al., 2022). The issue becomes more serious in an operational system because the input itself may no longer describe the current world.

1.2 Problem statement

Vehicle telemetry can be interrupted by tunnels and other signal-loss areas. During an interruption, a dispatch platform may keep the last received position and speed. The reading remains available, but its Age of Information (AoI) grows (Kaul, Yates and Gruteser, 2012). An attention map can therefore look complete even though some nodes represent old information. The interface reveals where the model placed attention, but it does not reveal whether that information is fresh or whether removing it would change the decision.

This creates an assurance problem. An operator may read a complete attention map as a trustworthy reason even when its evidence is stale, and stable dispatch performance cannot show whether the explanation is valid. This dissertation focuses on explanation reliability rather than increasing the number of pickups. It tests whether raw graph-attention weights provide reliable evidence about a decision under clean and degraded telemetry.

This dissertation uses *faithfulness decoupling* to describe a situation in which an explanation still appears credible after its supporting data have become stale, even though the displayed attention weights may not reliably identify the evidence that influenced the decision. One possible pattern is declining faithfulness while dispatch performance remains stable. The core question is whether attention provides reliable evidence under the tested conditions.

1.3 Research questions

The central research question is:

Can raw graph-attention weights be trusted as explanations of fleet-dispatch decisions when vehicle telemetry becomes stale?

Four operational questions provide the evidence needed to answer it:

1. **RQ1:** Under clean telemetry, do raw graph-attention weights identify decision-relevant nodes more reliably than type-matched random controls?
2. **RQ2:** Does tunnel-triggered telemetry degradation change the attention assigned to stale vehicle nodes relative to the paired clean observation?

3. **RQ3:** As outage duration increases, are changes in stale-node attention and decision-level faithfulness consistent across independently trained checkpoints?
4. **RQ4:** In the tested configurations, does degradation-aware training produce more consistent attention and faithfulness responses under telemetry degradation than clean training?

The construct-validity audit supports the interpretation of all four research questions. It is not a separate primary problem. Its purpose is to check that the faithfulness evaluator measures the effect of removing information rather than the accidental removal of an available action.

1.4 Objectives

The measurable objectives are to:

1. develop a reproducible Simulation of Urban Mobility (SUMO) fleet-dispatch benchmark that creates paired clean and degraded observations from the same simulated state and uses held-out test demand;
2. evaluate whether raw graph-attention weights identify decision-relevant nodes under clean telemetry using action-protected, type-matched controls and a leave-one-out (LOO) positive control;
3. quantify the paired change in attention assigned to stale vehicle nodes under tunnel-triggered outages and test its sensitivity to randomly triggered telemetry loss;
4. assess whether stale-node attention and decision-level faithfulness responses are consistent across outage durations, independently trained checkpoints, action groups, and attention-extraction choices;
5. compare the matched clean-trained and degradation-aware configurations to assess whether the tested degradation-aware configuration gives more consistent attention and faithfulness responses;
6. implement and apply a freshness-aware explanation audit framework that converts the collected evidence into an ELIGIBLE, WITHHOLD, or INCOMPLETE explanation-release decision.

1.5 Main findings

Overall, **this study does not validate raw graph-attention weights as reliable explanations of fleet-dispatch decisions**. Figure 1.1 summarizes the central finding: an attention explanation can remain available after the policy begins receiving stale vehicle information, but its availability does not show that it reliably identifies the evidence behind the decision.




Figure 1.1. Main finding. A tunnel-triggered outage leaves the policy with a stale observation. Attention remains available, but the audit withholds it because freshness and faithfulness are not both established.

Under clean telemetry, raw attention performs close to its type-matched random control and varies across the six trained policies. In contrast, the leave-one-out control produces a clearer faithfulness response. This shows that the evaluator can detect decision-relevant information, but raw attention does not consistently provide it.

Longer outages consistently increase exposure to stale vehicle information. However, the resulting changes in attention and faithfulness differ across trained policies. Some policies assign more attention to stale nodes, while others assign less. The study therefore does not find a single, consistent attention response to telemetry degradation.

The findings also depend on how attention is extracted and on whether the selected action is dispatch or no-op. Most scorable decisions are no-op, and the smaller dispatch-only results do not reproduce the main associations consistently. Degradation-aware training does not remove this variation.

The practical conclusion is that data freshness and explanation faithfulness must be checked separately for each trained policy. In the present experiment, the policy may still produce an action, but its raw attention weights should be withheld as an explanation.

1.6 Contributions

The dissertation makes four main contributions:

1. a reproducible paired clean/degraded benchmark that creates two observations from the same SUMO state, separates physical ground truth from policy input, and distinguishes the tunnel trigger from the observation-layer outage;
2. an action-aware, type-matched evaluation protocol that reduces the deletion bias caused by request nodes representing both information and available actions, checked through LOO, overlap, and action-type diagnostics;
3. evidence across independently trained checkpoints showing that stale-data exposure increases consistently with outage duration, while direct paired attention and DEF shifts vary across checkpoints and attention-extraction choices, with the pattern also tested under random telemetry loss;

4. an implemented freshness-aware explanation audit framework that combines freshness, faithfulness, action composition, extraction sensitivity, and checkpoint-consistency evidence to return an ELIGIBLE, WITHHOLD, or INCOMPLETE explanation-release decision.

1.7 Dissertation structure

Chapter 1: Introduction. Defines the research problem, research questions, objectives, main findings, and contributions.

Chapter 2: Literature Review. Reviews MARL fleet dispatch, graph attention, explanation faithfulness, graph explainability, and telemetry freshness.

Chapter 3: Methodology. Describes the SUMO environment, policy models, observation-layer degradation, faithfulness measures, evaluation design, and audit rules.

Chapter 4: Results. Reports policy capability, construct-validity checks, faithfulness results, stale-data exposure, sensitivity analyses, and audit decisions.

Chapter 5: Discussion. Interprets the findings, explains their limitations, and presents the freshness-aware explanation audit framework.

Chapter 6: Conclusion. Answers the research questions and summarizes the practical implications for trustworthy fleet-dispatch explanations.

Chapter 2: Literature Review

This chapter reviews the literature needed to define the research gap. It first considers reinforcement learning for relational fleet dispatch, then examines explanations for sequential and graph-based decisions. It also reviews the debate over attention and the limits of faithfulness tests. The final sections connect these topics to stale telemetry and position the experiment.

2.1 Reinforcement learning for fleet dispatch

Fleet dispatch links vehicles, passenger requests, and a road network over time. A decision for one taxi changes the demand available to others, so a collection of independent single-agent policies can create competition or duplicated assignments. Lin et al. (2018) model city-scale fleet management as a cooperative multi-agent reinforcement-learning (MARL) problem. Qin, Zhu and Ye (2022) show that reinforcement learning (RL) is used for matching, repositioning, pricing, and route decisions, often under changing supply and demand.

MARL offers several ways to represent coordination. Multi-Agent Deep Deterministic Policy Gradient (MADDPG) learns decentralized actors with centralized critics (Lowe et al., 2017). QMIX factorizes a team value into agent values while preserving a monotonic relation to the joint action (Rashid et al., 2020). Multi-Agent Proximal Policy Optimization (MAPPO) uses the more familiar Proximal Policy Optimization (PPO) objective with centralized training and decentralized execution; its empirical stability makes it a practical baseline for cooperative tasks (Yu et al., 2022). This dissertation adopts MAPPO because the research focus is the explanation attached to each taxi's local decision, not a new MARL algorithm.

Graphs provide a natural representation of the local dispatch problem. Nodes can represent the acting taxi, nearby taxis, and open requests; edges permit their information to be combined (Scarselli et al., 2009). A graph attention network (GAT) assigns learned weights while aggregating neighboring nodes (Velickovic et al., 2018). Actor-Attention-Critic similarly uses attention for communication in MARL (Iqbal and Sha, 2019). This makes GAT-MARL a suitable test case, not because it is assumed to be the best dispatch model, but because it matches the relational, multi-agent structure of fleet dispatch while exposing an attention channel that can be audited. These weights are easy to visualize, which makes them tempting to present as reasons for an action. Ease of display, however, is not evidence that the weights identify the information that caused the decision.

Recent dispatch methods continue to use relational structure. Hu, Feng and Li (2025) propose Localized Bipartite Match Graph Attention Q-learning (BMG-Q), which uses a local bipartite graph for ride-pooling dispatch. CoopRide studies cooperation across city grids (Wang et al., 2025), while DualG-MARL combines vehicle-state and task graphs (Sha et al., 2026). These studies evaluate dispatch or scheduling quality rather than whether their learned graph or attention weights can be presented as faithful explanations when vehicle telemetry becomes stale. They provide useful task benchmarks, but not direct faithfulness benchmarks.

2.2 Explainability in reinforcement learning

Explainable reinforcement learning (XRL) covers more than a visual saliency map. An explanation may describe which state features mattered, summarize a policy, contrast the

selected action with an alternative, or show the sequence of events that led to an outcome. Reviews by Heuillet, Couthouis and Diaz-Rodriguez (2021), Puiutta and Veith (2020), and Milani et al. (2024) all stress that the method must match the user's question. A model developer debugging a policy and an operator checking one live action do not need the same output. The systematic taxonomy by Bekkemoen (2024) reaches the same conclusion across a larger body of recent XRL studies.

This dissertation concerns a local, post-decision operator question: *which visible vehicles or requests supported the selected action, whether dispatch or no-op?* Earlier XRL work offers several ways to answer this question. Greydanus et al. (2018) perturb image regions to visualize what changes an Atari policy. Madumal et al. (2020) use a causal model to generate contrastive explanations. Mott et al. (2019) introduce an attention bottleneck intended to expose what an RL agent uses. These approaches differ in mechanism, but they share an important lesson: a useful picture and a faithful account of model dependence are separate properties.

For an operational dashboard, readability and faithfulness must be considered separately. An explanation may be easy to read but weakly connected to the policy computation. A faithful explanation should change when information that matters to the action is removed or retained. This experiment tests that link through a controlled decision-level faithfulness test; it does not measure user preference for the attention map.

2.3 Is attention an explanation?

A widely cited challenge to attention-based explanations was presented by Jain and Wallace (2019). They found that attention weights could be weakly correlated with other importance measures and that very different attention distributions could produce similar predictions. Wiegrefe and Pinter (2019) argued that this does not justify rejecting all attention explanations and proposed more careful tests. Serrano and Smith (2019) likewise showed that removing high-attention items can affect outputs, but that the relationship varies by model and task. The literature therefore supports neither a universal acceptance nor a universal rejection of attention. Each explanation claim needs an explicit definition and an appropriate test.

Attention became widely visible through Transformer models (Vaswani et al., 2017), but visibility also encouraged a broader interpretability claim than the mechanism alone can support. Lipton (2018) warns that the word *interpretability* often combines several different goals that should be evaluated separately.

Jacovi and Goldberg (2020) separate *plausibility*, which concerns whether an explanation looks reasonable to a person, from *faithfulness*, which concerns whether it reflects the model behavior that produced the output. This dissertation follows that distinction. It treats the attention weights as a candidate explanation and measures their decision relevance rather than accepting them simply because they are part of the actor.

Layer and head aggregation add another difficulty. A two-layer, multi-head network does not produce one unique attention map. Averaging heads, selecting a layer, taking a maximum, or composing attention across layers can produce different rankings. Abnar and Zuidema (2020) propose attention rollout and flow to account for information mixing across layers, and report stronger correlations with ablation and gradient importance than raw attention in Transformers. Although their model is not a GAT-MARL dispatcher, the methodological point transfers: an

attention result is incomplete unless the aggregation rule is defined and its sensitivity is checked.

2.4 Explaining graph neural networks

GNN explanations may identify important features, nodes, edges, or subgraphs. GNNExplainer learns a compact subgraph and feature mask that preserves a prediction (Ying et al., 2019). PGExplainer parameterizes explanation generation so that it can be shared across examples (Luo et al., 2020). SubgraphX searches for important subgraphs using Shapley-value approximations (Yuan et al., 2021). These methods show that graph explanation is a distinct problem: removing one node can alter both its own features and the messages available to other nodes.

The survey by Yuan et al. (2023) organizes GNN explainers by target, mechanism, and scope. GraphFramEx goes further by comparing explanation methods under different user needs and combining necessity and sufficiency views (Amara et al., 2022). Its results show that no single method dominates every evaluation dimension. This supports the use of several checks in the present study: DEF tests counterfactual relevance, taxi-only rank correlation compares attention with leave-one-out effects, and stale-attention measures describe freshness exposure.

Most GNN explainability benchmarks study node or graph classification. Fleet dispatch differs because some graph nodes also define actions. A passenger request is both information and a selectable action, while a nearby taxi is context only. Deleting them creates different interventions. This action-linked graph structure motivates the construct-validity audit in Section 3.7.

2.5 Evaluating faithfulness: methods and pitfalls

Comprehensiveness and sufficiency are common perturbation measures. In the ERASER benchmark, comprehensiveness asks how much a prediction weakens when the explanation is removed, while sufficiency asks how well the selected evidence alone preserves it (DeYoung et al., 2020). Liu et al. (2022) similarly test attention by measuring faithfulness violations. These measures are useful because they connect an explanation to model behavior instead of visual appeal.

Model-agnostic methods such as LIME also use local perturbations to estimate which inputs support a prediction (Ribeiro, Singh and Guestrin, 2016). Their usefulness depends on whether the perturbed samples remain meaningful for the model and task.

Perturbation is not automatically valid. Removing input features may create samples unlike the data seen during training. ROAR addresses this issue by removing features and retraining the model, showing why a simple deletion test can confound attribution quality with distribution shift (Hooker et al., 2019). Adebayo et al. (2018) provide a different sanity check: an explanation should respond when model parameters or training labels are randomized. Alvarez-Melis and Jaakkola (2018) show that explanation methods can also be locally unstable. Together, these studies show that the evaluator itself must be audited.

The present action-deletion problem is related but more specific. If a selected request node is masked, its action logit becomes unavailable. A large margin change can then occur even if the request features carried no useful evidence. A random baseline that removes more request nodes than the attention top-k is not a fair comparison. Type matching, selected-action

protection, and clamp logging are introduced to control this mechanism. A leave-one-out perturbation ranking acts as a positive control for comprehensiveness, while Gradient x Input provides a comparator that does not use attention coefficients.

The small graph creates a second limit. When $k = 3$ and only a few valid nodes exist, a random size-three set must overlap substantially with the explanation set. This compresses the possible difference between them. Reporting the exact expected overlap and a taxi-only rank correlation makes that limitation visible instead of treating every near-zero DEF as decisive evidence.

2.6 Age of Information and telemetry degradation

Age of Information (AoI) measures the time since the newest received update was generated (Kaul, Yates and Gruteser, 2012). It is widely used to study communication freshness, including the relationship between update age, estimation, and control (Yates et al., 2021). The explanation question is different. A dispatch action and its attention map may remain available even when a vehicle reading has stopped updating. The operator can see a strong weight without knowing that its source is old.

Delayed-observation reinforcement learning studies how an agent can continue to act when observations or actions arrive late. Bouteiller et al. (2021) model random action and observation delays and propose a delay-correcting actor-critic method. Liotet et al. (2022) learn a delayed policy from demonstrations produced without delay. Both studies focus on maintaining decision performance under delay. They do not test whether an explanation remains faithful when the policy receives stale telemetry.

This dissertation uses AoI only as a node-level freshness descriptor and as the weight inside weighted attention mass on stale nodes (WAMSN). The experiment manipulates the duration of an observation-layer outage triggered by tunnel entry, not AoI itself. It therefore makes no causal claim that a larger AoI value lowers faithfulness. The analysis instead asks whether stale exposure and attention-based explanation behavior provide consistent evidence under controlled outages.

2.7 From explanation evaluation to explanation assurance

An explanation test becomes an assurance process when its result is used to decide whether an explanation should be shown to a user. The NIST AI Risk Management Framework treats testing, validation, documentation, and monitoring as parts of managing the risks of AI systems (Tabassi, 2023). It provides general risk management guidance rather than a test for attention explanations.

This dissertation applies that general assurance principle to a narrower question: whether an attention map has enough supporting evidence to be presented for a fleet-dispatch decision. The resulting framework checks the data status, the validity of the faithfulness test, and consistency across independently trained policies. It then returns one of three decisions: *ELIGIBLE*, *WITHHOLD*, or *INCOMPLETE*. These labels and their decision rules are proposed in this dissertation; they are not prescribed by the NIST framework.

2.8 Research gap and positioning

The literature supports four points: relational RL is appropriate for fleet dispatch; attention is easy to expose but disputed as an explanation; GNN faithfulness requires graph-aware

perturbations; and stale state matters to operational control. Within the literature reviewed in this chapter, no study was identified that treats these topics as one explanation-assurance problem. This is a focused narrative review used to position the experiment; it is not a systematic or exhaustive literature search.

Recent graph-based dispatch studies provide the closest task comparison, but they answer a different question from this dissertation. Hu, Feng and Li (2025) evaluate BMG-Q, Wang et al. (2025) evaluate CoopRide, and Sha et al. (2026) evaluate DualG-MARL. Table 2.1 summarizes how their aims differ from the present study.

Table 2.1: Recent graph-based dispatch literature compared with this study

Study	Data and graph design	Main purpose	Unaddressed assurance question
BMG-Q (2025)	ride-pooling with a local bipartite matching graph	dispatch quality and scalability	Are graph weights faithful under stale telemetry?
CoopRide (2025)	cooperative MARL across city grids and real order data	city-scale dispatch quality	Are node-level explanations faithful?
DualG-MARL (2026)	vehicle-state and task graphs with real order data	ride-sharing scheduling quality	Do explanations remain supported under paired telemetry degradation?
This dissertation	local GAT-MAPPO taxi/request graphs in controlled SUMO runs	explanation assurance	Tests paired degradation, action-aware faithfulness, and release conditions.

These methods are not treated as performance baselines because their tasks, action spaces, and data-generation procedures differ. They are benchmark literature for positioning the contribution: recent work improves relational dispatch, whereas this dissertation tests whether an exposed attention channel has enough evidence to be presented as an explanation.

Within the graph- and RL-explainability and delayed-observation studies reviewed in Sections 2.2-2.6, no evaluation protocol was identified that combines all of the following:

- a graph-attention MARL action explained at decision level;
- paired clean and stale versions of the same operational observation;
- explicit measurement of attention placed on stale telemetry;
- a control for graph nodes that also remove available actions;
- replication across independently trained checkpoints.

This thesis contributes an audit protocol and an empirical test rather than a new dispatch algorithm. SUMO supplies a controlled simulated state (Lopez et al., 2018). The GAT-MAPPO policy supplies an auditable attention channel. The research contribution is to test whether that channel gives a reproducible freshness-aware and decision-relevant explanation, and to state clearly when the evidence cannot support that claim.

Chapter 3: Methodology

3.1 Study design

The study uses a controlled simulation because the research question requires the same decision to be observed in two ways: once with current telemetry and once with degraded telemetry. A conventional observational fleet log would not normally provide this exact clean twin. SUMO provides the simulated physical state, while a separate observation layer controls what the policy receives.

The protocol is defined in version-controlled configuration files. It contains three stages: training, validation-based checkpoint selection, and held-out evaluation. Faithfulness sweeps are run only after the selected checkpoints are frozen.

The six objectives are implemented through the paired observation benchmark, clean faithfulness controls, telemetry-degradation sweep, cross-checkpoint consistency analyses, matched training comparison, and final explanation-release audit.

3.2 SUMO environment and data

The road network is an OpenStreetMap extract of the Central Park area in Yubei, Chongqing. Roads marked as tunnels supply physically meaningful signal-loss triggers. SUMO simulates 20 taxis and 50 passenger requests during each 1,200-second episode. Taxi dispatch is controlled through TraCI, while SUMO continues to update the true position and speed of every vehicle.

Twenty generated demand variants use one fixed taxi seed (42) and passenger request seeds 1000-1019. Variants 0-13 form the training split, variants 14-16 form the validation split, and variants 17-19 form the held-out test split. The policy is fitted on the 14 training variants. Candidate checkpoints are compared on the three validation variants. The final reported runs use only the three test variants. This separation prevents test results from influencing model selection. During evaluation, consecutive episodes cycle through the three test variants, and actions are sampled from the policy's probability distribution. Evaluation seeds control this action sampling; they do not create additional independent demand datasets.

SUMO is a microscopic traffic simulator (Lopez et al., 2018), not a source of directly observed traffic data. Its outputs depend on the chosen road network, generated demand, routes, vehicle behavior parameters, tunnel locations, and random seeds. The experiment is therefore not described as free from bias. Instead, it uses these assumptions consistently in a paired design. At each sampled decision, the clean and degraded observations share the same underlying SUMO state and trained-policy checkpoint; only the observation-layer telemetry differs. This controls the comparison well enough to isolate the effect of the implemented degradation within this scenario, but it does not make the simulated fleet representative of every real city.

3.3 Observation graph and action space

Each idle taxi receives a self-centered graph. Table 3.1 summarizes its node types, features, and roles.

Table 3.1: Observation graph node types and roles

Node type	Main features	Role
Self taxi	position, normalized time, speed and AoI; separate position-valid indicator	acting vehicle
Peer taxi	relative position, availability, distance, AoI	nearby fleet context
Passenger request	pickup/drop-off vectors, waiting time	candidate request and action

The graph contains up to five nearby taxis and five nearby requests. Features are normalized, and request and taxi masks prevent padded nodes from receiving attention. The discrete action space contains no-op plus one action for each visible passenger request. Figure 3.1 shows the maximum graph size and the request-to-action mapping. This mapping is important for the faithfulness evaluation: masking a passenger request node also removes its associated action, whereas masking a peer-taxi node removes only information about that taxi. Section 3.7 explains how the evaluation accounts for the different effects of masking request and peer-taxi nodes.




Figure 3.1. Maximum configured graph: one self taxi, five peer taxis, and five requests. Valid nodes are fully connected, and unused slots are masked. The red ring marks stale telemetry; dark edges show illustrative attention into the self taxi. Each request maps to one dispatch action.

3.4 Policy models and training

The main policy uses per-node-type feature projections followed by two graph attention layers with four heads. The actor scores no-op and the visible requests. The critic estimates value for MAPPO training (Yu et al., 2022) with centralized training and decentralized execution (CTDE). The GAT implementation returns its attention tensors directly so they can be audited without changing the trained network.

During training, the centralized critic pools the valid node embeddings from all agents in the same transition and returns one shared team-value estimate. The actor does not receive this pooled representation: each taxi selects its action from its own local graph. At execution and during the explanation audit, the actor therefore remains decentralized.

The implementation uses scaled dot-product graph attention on the fully connected set of valid local nodes. This differs from the additive attention score in the original GAT paper, so "GAT"

in the experiment labels denotes the graph-attention policy family, not an exact reproduction of that layer. For node i , node j , layer l , and head h :

$$\begin{aligned}
q_i &= W_{Q,h} \text{LayerNorm}(h_i^l) \\
k_j &= W_{K,h} \text{LayerNorm}(h_j^l) \\
v_j &= W_{V,h} \text{LayerNorm}(h_j^l) \\
s_{ij}^{l,h} &= \frac{q_i k_j}{\sqrt{d_{\text{head}}}} \\
\alpha_{ij}^{l,h} &= \text{softmax}_j(s_{ij}^{l,h}) \\
z_i &= \text{concat}_h(\sum_j(\alpha_{ij}^{l,h} \times v_j)) \\
u_i &= h_i^l + W_O z_i \\
h_i^{l+1} &= u_i + \text{FFN}(\text{LayerNorm}(u_i))
\end{aligned} \tag{3.1}$$

Here, FFN denotes the feed-forward network within each attention layer. The mask removes padded nodes before the softmax. A residual connection keeps the previous node embedding. The actor reads the updated self embedding for no-op and each updated request embedding for its matching dispatch action. The default explanation is the self row ($i = 0$) averaged over both layers and all four heads, then renormalized:

$$\alpha_j = \text{normalize}\left(\frac{\sum_{l,h}(\alpha_{0j}^{l,h})}{LH}\right) \tag{3.2}$$

This aggregation is declared before evaluation. A sensitivity analysis also reports each layer, head-wise values, head maximums, and attention rollout; it does not select an aggregation after seeing which result is favorable. The explanation average is an analysis choice: inside the policy, head outputs are concatenated and projected rather than averaged.

The fixed self row represents the acting taxi's dashboard view and is the explanation channel originally declared for audit. It directly contributes to the no-op embedding. A request-action logit, however, is read from that request's updated embedding, so its corresponding request row is more directly aligned with the selected logit. A sensitivity analysis therefore reports a separate request-action sensitivity comparison between the fixed self row and the selected request row. This comparison is not used to redefine the primary metric after seeing the result.

Figure 3.2 summarizes the policy path and the read-only audit channel.

Graph-Attention Policy and Read-Only Audit




Figure 3.2. The policy and audit share the same two-layer attention encoder. The default audit reduces the self-node attention row to node importance without modifying the actor, critic, or selected action; the selected-request row is evaluated separately as a sensitivity check.

Table 3.2 lists the three policy conditions included in the experiment. The non-graph baseline uses a multilayer perceptron (MLP).

Table 3.2: Policy conditions used in the experiment

Model	Policy	Training observations	Purpose
MLP	MLP-MAPPO	clean	performance context without graph attention
GAT	GAT-MAPPO	clean	primary attention model under audit
GAT-Outage	GAT-MAPPO	tunnel-triggered 30-second outages	tested degradation-aware training configuration

Each model is trained in three independently initialized runs using seeds 42, 43, and 44. Corresponding GAT and GAT-Outage seeds are paired in the H5 comparison.

The maximum training budgets were fixed from validation-only stability diagnostics before the held-out evaluation. Clean GAT uses 40 epochs because longer development runs showed late policy drift. MLP uses 50 epochs because it learned too slowly under the lower GAT learning rate, while a 60-epoch probe failed the 80% final-validation-retention gate. GAT-Outage uses 50 epochs because its seed-43 validation result was still improving at epoch 40 and retained 93.9% of its selected validation pickup count at epoch 50.

These choices improve stability but mean that GAT and GAT-Outage do not have identical optimization horizons: their learning rates decay over 40 and 50 epochs, respectively. H5 therefore compares the two fitted training configurations descriptively. It does not isolate the effect of degraded training observations from the difference in training budget.

Candidate checkpoints include snapshots saved at zero-based indices 0, 10, 20, 30, and so on, the final snapshot, and a rolling-best snapshot updated when the ten-epoch mean pickup count improves. Because the implementation records zero-based indices, index 39 denotes the checkpoint saved after 40 training epochs. Selection uses eight sampled-action validation

episodes with seed 2026 and mean pickups as the primary criterion; mean reward and the earlier checkpoint index break ties. The test split is not read during selection. Table 3.3 reports the selected checkpoint index for each training run.

Table 3.3: Validation-selected checkpoint indices (zero-based)

Model	seed 42	seed 43	seed 44
MLP	20	40	34
GAT	39	39	30
GAT-Outage	49	40	40

Table 3.4 reports the training and optimization settings.

Table 3.4: Training and optimization settings

Setting	Value
Training epochs	40 (GAT), 50 (MLP and GAT-Outage)
Learning rate	0.0001 (GAT variants), 0.0003 (MLP)
Learning-rate schedule	linear decay over each run (GAT variants); constant (MLP)
Discount factor (gamma)	0.99
Generalized advantage estimation factor (lambda)	0.95
PPO clip ratio	0.20
PPO passes per update	4
Minibatch size	256
Value-loss coefficient	0.25
Value clip ratio	0.20
Entropy coefficient	0.05, adaptive to a 0.20 entropy floor
Maximum adaptive entropy coefficient	0.20
Target Kullback-Leibler divergence	0.015
Gradient-norm limit	0.50
Dispatch credit reward	1.00
Critic encoder gradient scale	0.10
GAT hidden width / layers / heads	64 / 2 / 4
GAT attention dropout	0
MLP hidden width / layers	176 / 3

Table 3.5 separates checkpoint acceptance from the optimizer settings. Final validation retention is the final checkpoint's mean validation pickups divided by the selected checkpoint's mean validation pickups. The stability window defines where the training measures are calculated; the remaining four gates must pass before a training run is accepted for evaluation.

Table 3.5: Checkpoint-acceptance criteria

Checkpoint-acceptance setting	Required value
Stability window	Final 10 training epochs
Minimum final-window mean pickups	5.0
Maximum consecutive zero-pickup epochs	4
Minimum selected validation pickups	5.0
Minimum final validation retention	0.80

At simulation step t , the reported shared team reward is:

$$r_t = 10 N_{\text{pickup},t} + 0.5 N_{\text{dispatch},t} - 0.001 \text{mean_wait}_t \quad (3.3)$$

For PPO training, taxi i receives an additional difference credit only when its dispatch is accepted:

$$r_{i,t}^{\text{train}} = r_t + 1.0 I_{\text{successful dispatch by } i \text{ at } t} \quad (3.4)$$

Here, $I(\cdot)$ is one when the condition is true and zero otherwise. The shared reward is retained for reporting, while the agent-specific training reward enters advantage estimation and PPO updates. Neither reward contains a term for a faithful, stable, or human-readable attention map. Figure 3.3 shows how training, validation selection, and held-out evaluation remain separated.

Model Training and Leakage-Controlled Evaluation




Figure 3.3. Shared training and validation-based selection pipeline with the model-specific budgets described above. Each frozen checkpoint receives 24 clean held-out episodes. The faithfulness audit then compares the frozen GAT and GAT-Outage checkpoints across clean and four outage conditions without reselection. This produces 720 faithfulness-sweep episodes per GAT model and 1,440 in total.

3.5 Telemetry degradation

The trigger and the degradation location are separate parts of the design.

Step 1 - Trigger: when a taxi enters a tunnel edge, the event starts one outage.

Step 2 - Observation-layer outage: for the declared duration, the policy sees the taxi's last valid position and speed. SUMO continues to simulate the true state.

Step 3 - Recovery: after the fixed window expires, a current reading is accepted. A taxi must leave and enter a tunnel again before another tunnel-entry event can trigger.

The fixed outage durations are 10, 20, 30, and 60 seconds. They are easy to monitor and compare, and they create increasing empirical degradation rates. They are not treated as direct causal doses of explanation faithfulness.

AoI is calculated as the current simulation time minus the time of the last valid update and is included in each GAT observation. WAMSN uses normalized AoI as a staleness weight. For

this reason, an increase in WAMSN with outage duration partly verifies that the manipulation created longer stale exposure; it is not by itself proof that AoI changed DEF. Figure 3.4 illustrates this separation between physical state and policy observation.




Figure 3.4. Tunnel entry triggers an observation-layer outage. SUMO keeps moving the vehicle while the policy sees its last valid position; the clean twin reads current state but never acts. The paired observations keep the same node identities, node order, masks, and selected action.

3.6 Explanation measures

3.6.1 Decision-level explanation faithfulness

DEF asks whether the nodes ranked highly by an explanation affect the chosen action more than a size- and type-matched random set. Two score functions are recorded from the same counterfactual forward passes. **Probability DEF** uses the selected action probability, $p(G,a)$, and is bounded by the probability scale. **Logit-margin DEF** uses the selected action logit minus the strongest available alternative, $m(G,a)$. The margin can retain resolution when the softmax probability is saturated, but its value depends on a checkpoint's logit scale.

For top- k nodes R_k , comprehensiveness measures the output change when those nodes are removed; sufficiency measures the output retained when only those nodes remain. The self node is always retained and is not a deletion candidate. The candidate set contains only valid peer-taxi and request nodes. The two gains are compared with five matched random subsets for each k in $\{1,2,3\}$ and then averaged.

Let $s(G,a)$ denote either $p(G,a)$ or $m(G,a)$. The action a is sampled from the policy observation being audited. For a degraded/clean pair, the action sampled on the degraded side is held fixed when both observations and every counterfactual subset are scored. For explanation subset R_k and each matched random subset indexed by b :

$$\begin{aligned}
\text{Comp}(R_k, a) &= s(G, a) - s(G \setminus R_k, a) \\
\text{Suff}(R_k, a) &= s(G, a) - s(G[R_k], a) \\
g_{\text{comp}}(k; a) &= \text{Comp}(R_k, a) - \text{mean}_b \text{Comp}(B_k^b, a) \\
g_{\text{suff}}(k; a) &= \text{mean}_b \text{Suff}(B_k^b, a) - \text{Suff}(R_k, a) \\
\text{DEF}(a) &= \text{mean}_k \frac{1}{2} [g_{\text{comp}}(k; a) + g_{\text{suff}}(k; a)]
\end{aligned} \tag{3.5}$$

Comprehensiveness is better when removing the explanation weakens the action more than removing random nodes. Sufficiency is better when keeping the explanation loses less evidence than keeping random nodes. The logit margin is clipped to $[-10, 10]$ only when an action becomes unavailable or unopposed; every clamp is logged.

Type-matched random subsets are the **control baseline**; they are not themselves a faithfulness score. Corrected DEF is the measured difference between the attention-ranked subset and this matched random control. Its sign is interpreted as follows:

1. **Positive DEF:** the attention-ranked nodes affect the selected action more than comparable random nodes. This supports decision-level faithfulness.
2. **DEF near zero:** the attention ranking has no measured advantage over the matched random control.
3. **Negative DEF:** the attention-ranked nodes affect the selected action less than the matched random nodes. This does not support the attention ranking as a faithful explanation.

These interpretations apply to both variants. The predeclared H1-H4 analysis family and the paired clean-twin comparison use probability DEF, following the declared analysis plan. Logit-margin DEF is used for the construct-validity and ranking-control diagnostics because it can show output movement that a saturated probability hides. It is action-protected where request deletion could remove the chosen action. The two variants have different units and cannot be converted into one another. Both use the type-matched baseline described in Section 3.7. DEF provides comparative perturbation evidence; even a positive value does not prove that attention is a complete causal explanation.

The study uses related measures for different questions. They should not be read as interchangeable estimates. Table 3.6 states the role of each measure.

Table 3.6: Explanation measures and roles

Measure	Decisions and score	Role in the study
Probability DEF	all scorable sampled actions; probability scale	primary H1-H4 and paired clean/degraded analysis
Logit-margin DEF	all scorable sampled actions; selected-action margin	construct-validity and ranking diagnostics when probability is saturated
Dispatch self-row margin-DEF	sampled dispatch actions only; declared self-row explanation	release check for decision relevance under clean and 60-second outage conditions
WAMSN	stale-exposed decisions; AoI-weighted attention	freshness-exposure indicator and H3 manipulation check
Paired attention and DEF shifts	degraded observation minus its clean twin	direct within-decision response to the observation-layer intervention

3.6.2 Stale-node attention

WAMSN measures attention attached to stale vehicle information:

$$\text{WAMSN} = \frac{\sum_i (\text{attention}_i \times \text{normalized_AoI}_i)}{\sum_i (\text{attention}_i)} \quad (3.6)$$

Only self and peer-taxi nodes contribute because request nodes do not carry vehicle telemetry. The analysis reports WAMSN when at least one stale vehicle node is visible, together with stale-node attention share, stale-node mass, and whether a stale node enters the top three.

The paired stale-attention shift compares the degraded observation with its exact clean twin at the same simulation step. Positive values mean the degraded observation assigns more attention mass to nodes marked stale; negative values mean it assigns less.

3.7 Construct-validity audit

The audit supports the primary research problem by checking whether DEF is a valid comparison in this action-candidate graph.

A uniform random occlusion can select passenger requests more often than the attention top-k. Removing such a request changes both the observation and the available action set. A nearby-taxi occlusion changes only information. The result can therefore reflect different node-type composition rather than explanation quality.

The corrected protocol uses three controls:

- **type matching:** random subsets contain the same mixture of taxi and request nodes as the explanation subset;
- **chosen-action protection:** the request associated with the selected action is not deleted in the action-protected variant;
- **clamp logging:** counterfactual action-margin clamps are counted so that action deletion can be detected.

Figure 3.5 illustrates why these controls are required.

Why Node Occlusion Needs Construct-Validity Controls




Figure 3.5. Request-node occlusion can remove its action, while peer-taxi occlusion removes context only. Corrected DEF matches subset size and node types, protects the chosen request, and logs margin clamps.

Uniform-baseline results are retained only as an audit trail. They are not used for the final hypothesis verdicts.

Three additional diagnostics test whether a near-zero DEF can be interpreted. First, a leave-one-out (LOO) control masks each valid node separately and ranks nodes by the resulting selected-action margin loss. The selected request is protected. This control is expected to strengthen comprehensiveness because it is built from the same single-node perturbation, but it is not an independent proof that the combined DEF is optimal. Second, Gradient x Input (GxI) ranks each node by the L2 norm of the selected-action logit gradient multiplied by its raw features. It is a post-hoc comparator that does not use attention weights. Third, the audit reports valid taxi/request counts and the exact expected top-k overlap with a type-matched random subset. Taxi-only Spearman correlation between attention and LOO effects supplies a ranking measure that does not delete request actions and does not depend on random subset overlap.

3.8 Evaluation design

3.8.1 Primary evaluation

MLP, GAT, and GAT-Outage are evaluated under clean telemetry with eight evaluation seeds (42-49) and three episodes per seed. One checkpoint is frozen for each model and training seed, giving 24 held-out episodes per checkpoint. The clean capability evaluation contains 216 episode evaluations:

$$3 \times 3 \times 8 \times 3 = 216 \quad (3.7)$$

Table 3.7 summarizes the complete evaluation structure. An episode evaluation is one complete rollout of one frozen checkpoint under one telemetry condition.

Table 3.7: Held-out evaluation structure and episode counts

Evaluation	Model families	Total frozen checkpoints	Telemetry cases	Episodes/cell	Total
Clean capability context	3	9	1	24	216
Exposure-conditioned faithfulness sweep	2	6	5	48	1,440
Added ranking controls	2	6	2	24	288
Random-trigger sensitivity analysis	2	6	3	24	432

GAT and GAT-Outage receive the full faithfulness sweep:

$$2 \times 3 \times 5 \times 8 \times 6 = 1,440 \quad (3.8)$$

This gives 720 episode evaluations per GAT model and 1,440 across both models.

Table 3.8 lists the five telemetry conditions and their roles.

Table 3.8: Telemetry conditions used in held-out evaluation

Telemetry condition	Observation-layer treatment	GAT-Outage relation	Evaluation purpose
Clean	Current vehicle readings	No imposed outage	Capability context and faithfulness baseline
10 s	Freeze last valid reading for 10 seconds	Unseen shorter duration	Mild-degradation transfer test
20 s	Freeze last valid reading for 20 seconds	Unseen shorter duration	Intermediate transfer test
30 s	Freeze last valid reading for 30 seconds	Training-matched duration	In-condition degradation test
60 s	Freeze last valid reading for 60 seconds	Unseen longer duration	Severe stress and transfer test

The legal-random capability baseline samples uniformly from no-op and the currently valid request actions. The greedy baseline is deliberately naive: each idle taxi chooses its nearest visible request independently. If several taxis choose the same request, the environment accepts the first action in the step and rejects the later conflicts; there is no fleet-wide assignment or conflict optimization. The policy is deterministic for a fixed demand variant, so its uncertainty is based on three independent demand variants rather than the repeated evaluation-seed records. It is therefore interpreted only as a lower bound, not as a competitive dispatch algorithm.

The relation column refers only to GAT-Outage; the clean-trained GAT has not seen any imposed outage duration during training. The evaluation changes the observation-layer condition but never retrains or reselects a checkpoint. Outside stale exposure, faithfulness is sampled every 16 decisions. Every decision that contains at least one stale vehicle node is scored. Each such observation is also compared with its exact clean twin at the same simulation step. The pair keeps the same trained checkpoint, agent, graph-node identities, node order, validity masks, and sampled action. Only telemetry-dependent feature values and freshness indicators differ. The degraded and clean evaluations also reuse the same matched-random subsets, so the paired difference does not contain avoidable baseline-sampling noise.

Each degraded condition must contain at least 20 stale-exposed records from at least three episodes. Preflight also checks expected cells, clean source provenance, type-matched controls,

chosen-action exclusion, complete event capture, and separation in empirical AoI between conditions. A checkpoint sweep is included in the analysis only if it passes these gates.

3.8.2 Supporting analyses

Actions in the primary learned-policy evaluation are sampled from the policy distribution, matching training and preserving both no-op and dispatch decisions for audit. A separate descriptive diagnostic applies deterministic argmax to each of the six GAT checkpoints over three held-out episodes. It does not affect checkpoint selection or the primary capability comparison.

The action-type diagnostic separates scorable decisions into no-op and dispatch; cases with no available request are excluded because no-op is then forced. Records are averaged by episode within each action group. This shows whether combined results apply to both action types or mainly reflect the more common action.

The random-trigger sensitivity analysis evaluates the six frozen checkpoints under clean telemetry, a 30-second tunnel-triggered freeze, and a 30-second randomly triggered freeze. A validation-only calibration set the per-taxi, per-step trigger probability to 0.0023, giving 0.569% exposure for the reference checkpoint compared with 0.580% under the tunnel trigger. The probability is then fixed for every checkpoint. Realized exposure is reported because actions change taxi availability and trajectories. This post-study analysis tests transfer to another trigger mechanism; it is not a fully independent robustness test because the calibration used a reference checkpoint and was not preregistered.

The ranking controls use the same held-out demand and action-sampling protocol. They evaluate GAT and GAT-Outage under clean telemetry and a 60-second outage:

$$2 \times 3 \times 2 \times 8 \times 3 = 288 \quad (3.9)$$

LOO, Gradient x Input, overlap, and aggregation diagnostics are sampled every eight decisions. Query-row sensitivity is evaluated for every selected-request action. Records are averaged within episodes before bootstrap intervals are calculated. Attention-aggregation sensitivity is reported by training seed for decisions with at least one visible stale vehicle node.

3.9 Hypotheses and statistical analysis

Table 3.9 provides one map of the five hypotheses, their analyses, and their roles. H1, H3, and H4 were predeclared as primary tests; H2 was predeclared as exploratory; and H5 is a matched-seed descriptive comparison.

Table 3.9: Hypothesis map and analysis roles

Hypothesis	Question	Analysis	Role
H1	DEF across outage durations	episode-block probability-DEF trend	primary
H2	Faithfulness versus capability decline	paired clean-relative change	exploratory
H3	WAMSN across outage durations	episode-block WAMSN trend	primary check
H4	Within-episode WAMSN-DEF link	Spearman correlation and sign-flip test	primary
H5	Weaker relationships for GAT-Outage	matched-seed correlation comparison	descriptive

RQ1 uses clean DEF, type-matched random controls, LOO, and ranking controls. RQ2 uses the paired clean/degraded stale-attention comparison. RQ3 uses H1, H3, and H4 together with the paired and action-type diagnostics. RQ4 uses H5.

3.9.1 Statistical procedure

H1 and H3 use episode-block permutation trend tests and cluster bootstrap confidence intervals. H2 uses paired cell-level sign flips relative to each evaluation seed's clean condition. H4 first calculates a Spearman correlation within each episode and then applies a sign-flip test. Holm adjustment is applied to H1-H4 within each trained policy; H2 remains exploratory.

All tests use 10,000 permutation draws, 2,000 bootstrap draws, statistical seed 0, and 95% percentile bootstrap intervals. Episodes, rather than individual decisions, are the resampling blocks. Each training seed represents one trained policy. Three seeds can show variation but cannot support a population-level claim, so cross-seed results are reported as ranges and counts of supporting seeds without a population p-value.

H5 is supported for a matched seed only when both the outage-duration/DEF and WAMSN/DEF correlations are closer to zero for GAT-Outage than for GAT. The result is the number of supporting pairs out of three; no cross-seed significance test is used.

3.9.2 Interpretation limits

H2's original rate divides by clean DEF. Because clean DEF is close to zero, that ratio can exaggerate a small absolute change. The analysis therefore uses the support verdict and absolute trajectories as its main evidence.

H3 is a manipulation check, not evidence that AoI causes DEF to change. WAMSN contains normalized AoI by definition, so H3 only confirms that longer outages create more AoI-weighted stale exposure.

H5 compares two fitted configurations with different maximum training budgets and learning-rate decay horizons. It does not isolate the effect of degraded training observations. The exposure-conditioned paired follow-up is also kept separate from H1-H5: it subtracts clean-twin DEF from degraded DEF for the same decision and reports episode-block intervals, effect size, direction across trained policies, and uncertainty.

3.10 Audit decision implementation

The experiment pipeline produces measurements; a separate release-audit stage turns those measurements into a decision about whether attention may be shown as an explanation. The stage reads saved evidence and leaves the policy, selected action, SUMO state, and observations unchanged.

Its inputs are the frozen-checkpoint preflight reports, condition summaries, action-type results, training-seed synthesis, deterministic diagnostics, faithfulness controls, and tunnel-versus-random trigger comparison. These artifacts jointly cover telemetry freshness, decision relevance, action composition, attention extraction, checkpoint replication, and the action rule intended for deployment.

The implementation applies nine checks:

1. evidence integrity;
2. separate freshness reporting;
3. type-matched and action-protected controls;
4. dispatch-action decision relevance;

5. no-op and dispatch composition;
6. attention-extraction stability;
7. consistency across independently trained checkpoints;
8. capability of the action rule intended for deployment;
9. robustness to tunnel and random loss triggers.

Decision relevance passes only when the 95% confidence-interval lower bound of self-row dispatch margin-DEF is greater than zero under both clean and 60-second outage evaluation. A positive value means that the attention ranking is more informative than its type-matched random control. Extraction stability fails if an audited layer, head, maximum, or rollout rule reverses the direction produced by the declared self-row mean. Trigger robustness uses confidence intervals: a direction is supported only when its 95% interval excludes zero, and supported tunnel and random-trigger directions must agree. An interval that crosses zero is INDETERMINATE, not evidence of a reversal. The remaining rules test evidence availability, action coverage, checkpoint agreement, and the action rule declared for deployment.

The present audit concerns actions sampled from the policy distribution. The argmax evaluation is therefore retained as a descriptive limitation rather than a mandatory release gate. If an argmax policy were intended for deployment, its held-out capability check would become mandatory. The current action-composition gate requires both action groups and at least 20 dispatch decisions at the model and individual-checkpoint levels, matching the minimum stale-exposure record floor used by the evaluation preflight. It confirms that dispatch evidence is present but does not by itself establish that a sparse dispatch group is representative.

The output has three states. ELIGIBLE means that every required check passed within the stated scope. WITHHOLD means that at least one check failed and attention must remain an internal diagnostic. INCOMPLETE means that required evidence is missing, so no release decision can be made. Eligibility is not proof of a complete causal explanation; it only permits presentation with an explicit freshness indicator and audit scope as an audited candidate explanation.

The release rules are a post-study synthesis of the risks identified by the experiment. Their application to the present results is therefore demonstrative and is not an additional confirmatory test of H1-H5. For future candidate models, the rules and thresholds must be fixed before held-out evaluation.

3.11 Reproducibility

The experiment runner records the configuration, commands, seeds, checkpoint hashes, source revision, and preflight reports in versioned run directories. Compact outputs include the summary, performance context, deterministic diagnostic, action-type audit, and training-seed synthesis tables. The training-seed synthesis makes consistency across trained policies explicit. Reproduction commands, file locations, and expected outputs are documented in docs/REPRODUCE_EXPERIMENTS.md. The complete frozen-checkpoint workflow is available as a single framework stage. It produces a machine-readable JSON decision, a CSV check table, and a reader-facing Markdown report. The detailed input and output contract is documented in docs/FRESHNESS_AWARE_EXPLANATION_AUDIT.md.

3.12 Ethics and data governance

The experiment uses a simulated taxi fleet, generated demand, and an OpenStreetMap-derived road network. It contains no human participants, personal records, live vehicle identifiers, or commercial dispatch data. Its main ethical risk is miscommunication: presenting attention as a trustworthy reason could give an operator false confidence. The reporting therefore keeps freshness, task behavior, and faithfulness separate and avoids unsupported causal claims.

Chapter 4: Results

The automated preflight audit passed 100 checks covering provenance and protocol requirements, and all six GAT checkpoint sweeps passed the inclusion gates defined in Section 3.8. Results are reported by training seed because decisions and episodes from one checkpoint do not constitute independent model replications. Unless stated otherwise, uncertainty is reported as a 95% confidence interval (CI).

4.1 Policy capability and training stability

Validation selected GAT checkpoint indices 39, 39, and 30 and GAT-Outage indices 49, 40, and 40 for seeds 42, 43, and 44. These are zero-based indices; for example, index 39 is the checkpoint saved after 40 training epochs. The selected index is reported separately for each trained policy.

Figure 4.1 shows the training and checkpoint-selection diagnostics.




Figure 4.1. Training diagnostics for GAT (upper panels) and GAT-Outage (lower panels). Curves use a ten-epoch moving mean; stars mark validation-selected checkpoints. Entropy remains above the collapse threshold, and PPO clipping is limited.

All nine training runs pass the declared stability gates. The following values describe the final ten training epochs, not the held-out evaluation in Table 4.1. For GAT, the final ten-epoch mean pickups are 13.3, 9.0, and 9.9, and final validation retention is 100%, 100%, and 92.6%. The corresponding GAT-Outage values are 12.4, 14.1, and 13.8 pickups, with validation retention

of 100%, 93.9%, and 96.5%. Training metrics still fluctuate across epochs, but no selected checkpoint is a zero-pickup policy under sampled-action validation. Deterministic argmax results are reported separately below.

Table 4.1 places the selected policies above two lower bounds evaluated on the same held-out demand. Performance is reported only to establish that the audited policies perform above random and greedy lower bounds.

Table 4.1: Policy capability on held-out demand

Policy	Mean pickups	Uncertainty or replicate range
Legal random	1.50	95% episode-cluster CI [1.08, 1.96]
Greedy nearest request	2.00	95% demand-variant CI [1.00, 3.00]
MLP	13.63	training-seed range 12.83-14.17
GAT	11.64	training-seed range 9.42-13.71
GAT-Outage	12.97	training-seed range 11.79-14.58

The legal-random baseline has 24 independent episode clusters. Greedy nearest is deterministic, so its effective sample is the three demand variants rather than 24 repeated seed-episode records. GAT's seed means are 13.71, 9.42, and 11.79. Every selected policy exceeds both lower bounds when actions are sampled from its probability distribution. These results confirm capability for the sampled policy distributions only; the role of graph attention and the usability of the argmax policy remain untested.

The audit requires each checkpoint's sampled policy distribution to exceed the declared lower bounds, not GAT to outperform MLP. That condition is met. MLP has the highest cross-seed mean in this implementation, although the replicate ranges overlap. The experiment finds no task-performance advantage for graph attention. This limits how broadly the audit can be generalized but remains separate from whether the exposed weights faithfully explain GAT decisions.

Figure 4.2 shows the held-out pickup count for each selected policy.




Figure 4.2. Each point is one selected policy evaluated over 24 held-out episodes. Horizontal lines show the cross-seed mean.

The deterministic diagnostic changes the capability interpretation. All six selected graph-attention checkpoints choose no-op throughout three held-out episodes, giving 0 pickups in 18 of 18 checkpoint-episodes. Their common mean reward is -46.58 and final pending-request wait is 905.9 seconds. The sampled-action results show only that the policy distributions assign enough probability to dispatch actions to produce pickups when sampled; they do not establish a stable argmax dispatcher.

4.2 Construct-validity audit

Removing a request node can remove an available action, while removing a peer taxi only hides information. A uniform random occlusion therefore compares different types of change. Table 4.2 quantifies the resulting control-baseline artifact.

Table 4.2: Construct-validity audit of random controls

Control baseline used to calculate DEF	Mean margin-DEF	95% CI	Interpretation
Uniform random	-0.5540	[-0.5808, -0.5288]	dominated by node-type and action-deletion imbalance
Type-matched random	-0.0093	[-0.0159, -0.0013]	small negative result after matching node types

The type-matched interval excludes zero, but its magnitude is about two orders smaller than the uniform-control artifact. This 1,199-decision audit supports the measurement design rather than answering the primary research question. The reported protocol uses type matching, protects the chosen request in the margin variant, and logs action-margin clamps.

4.3 Evaluator sensitivity and ranking controls

Raw attention, Gradient x Input, and an LOO perturbation ranking are evaluated under clean and 60-second conditions. Each model, training seed, and condition contains eight evaluation seeds and three episodes. Table 4.3 shows clean results. The 60-second aggregate means remain close to clean.

The two conditions are stored as separate evaluations and have different decision counts. Across the six checkpoints, the largest absolute 60-second minus clean change is 0.000071 for raw attention, 0.000254 for Gradient x Input, and 0.000306 for LOO. Figure 4.3 shows the clean values and the condition difference separately instead of presenting two visually indistinguishable rows.

Table 4.3: Clean-telemetry faithfulness controls by training seed

Model	Seed	Raw-attention margin-DEF	GxI margin-DEF	LOO margin-DEF	Taxi-only rho
GAT	42	+0.0003	-0.0055	+0.0008	+0.088
GAT	43	-0.0009	-0.0033	+0.0099	+0.404
GAT	44	+0.0015	+0.0012	+0.0030	+0.426
GAT-Outage	42	+0.0011	-0.0019	+0.0020	+0.472
GAT-Outage	43	+0.0006	-0.0010	+0.0089	+0.602

Model	Seed	Raw-attention margin-DEF	GxI margin-DEF	LOO margin-DEF	Taxi-only rho
GAT-Outage	44	+0.0031	+0.0052	+0.0044	+0.699

The LOO-ranked control is positive and larger than raw attention for all six selected checkpoints. This shows that DEF responds to a ranking built from single-node margin loss. LOO is still a perturbation control, not a ground-truth explanation. Gradient x Input changes sign across checkpoints and provides no uniform improvement.

Raw-attention DEF ranges from -0.0009 to +0.0031. Taxi-only rank correlation with LOO is positive in all six checkpoints, but ranges from 0.088 to 0.699. Attention shows some agreement with decision-relevant LOO ordering, but the strength of that agreement remains checkpoint-dependent.

Faithfulness controls and the independently evaluated 60 s change




Figure 4.3. Top: clean logit-margin DEF. Bottom: 60-second minus clean DEF, multiplied by 10,000, from separate evaluations. Lines join methods for each checkpoint; non-zero changes confirm distinct records.

The attention reduction rule remains important. Across the declared mean, individual heads, layer maxima, and rollout, every checkpoint contains both a positive and a negative stale-attention shift. The widest ranges are -0.0109 to +0.0063 for GAT seed 44 and -0.0097 to +0.0064 for GAT-Outage seed 44.

Figure 4.4 shows the sensitivity to the attention reduction rule.




Figure 4.4. Stale-attention shift multiplied by 1,000. Selecting a layer, head, or rollout rule after seeing the result could reverse the conclusion.

Changing from the declared self row to the selected request row has little effect for seeds 42 and 44, but increases request-action DEF by about 0.0064 for GAT seed 43 and 0.0073 for GAT-Outage seed 43. The aggregate gap is +0.0022 for GAT and +0.0025 for GAT-Outage because it averages these different policy responses. The largest absolute 60-second minus clean query-row change is 0.000264, so the two conditions again remain close without being duplicates.

Figure 4.5 compares the declared self row with the selected request row.

Query-row sensitivity and the independently evaluated 60 s change




Figure 4.5. Top: clean request-action DEF from the declared self row and the selected request row. Bottom: each query row's independently evaluated 60-second result minus clean, multiplied by 10,000. The selected request row improves the seed-43 clean results but does not provide a consistent improvement across checkpoints or conditions.

4.4 Small-graph resolution

The scored graphs usually contain the maximum number of visible nodes, as summarized in Table 4.4.

Table 4.4: Valid node-count distribution in scored graphs

Model	Node count	Mean	Median	5th percentile	95th percentile
GAT	peer taxis	5.00	5	5	5
GAT	requests	4.51	5	1	5
GAT	non-self total	9.51	10	6	10
GAT-Outage	peer taxis	5.00	5	5	5
GAT-Outage	requests	4.52	5	1	5
GAT-Outage	non-self total	9.52	10	6	10

Even so, type matching creates substantial top-k overlap. The expected random intersection is about 0.22 for $k=1$, 0.41 for $k=2$, and 0.60 for $k=3$. Attention-LOO agreement at $k=1$ exceeds random overlap for both models. At $k=3$, it is 0.60 for GAT and 0.49 for GAT-Outage, compared with expected overlap of about 0.60. Restricting analysis to at least six non-self nodes does not change the sample because every scored decision already meets that threshold. This pattern suggests that the ranking agreement has more contrast at $k=1$; it does not establish that LOO is ground truth or that the top-ranked attention node is a causal explanation.

Figure 4.6 shows how the expected overlap grows with k .




Figure 4.6. Large expected overlap between the explanation and its matched random set reduces the contrast available to DEF, especially at $k=3$.

4.5 Telemetry manipulation and stale exposure

Tunnel-triggered observation outages create increasing stale exposure. The event-aware audit scores every decision that contains at least one stale taxi, instead of relying on periodic sampling. Across GAT checkpoints, the number of stale-exposed records ranges from 347-382 at 10 seconds and 1,476-1,806 at 60 seconds. The GAT-Outage ranges are 353-372 and 1,566-1,697. Every degraded cell exceeds the pre-declared coverage gates.

Conditional WAMSN rises from 0.0275-0.0289 at 10 seconds to 0.0804-0.0900 at 60 seconds. H3 is supported for all six selected checkpoints. Because WAMSN contains normalized AoI, this verifies increasing stale exposure; it does not show that AoI causes lower faithfulness. Figure 4.7 displays the duration sweep.




Figure 4.7. Longer observation outages increase WAMSN for every policy. DEF changes are much smaller; pickups are reported only as context, not as the main research outcome.

4.6 Exposure-conditioned paired audit

The paired audit compares each degraded observation with its exact clean twin at the same simulation step. Positive attention shift means more attention is assigned to nodes marked stale. Negative DEF shift means lower measured faithfulness under degradation. Table 4.5 reports the paired estimates.

Table 4.5: Exposure-conditioned paired attention and DEF shifts

Model	Seed	Attention shift x 1,000 [95% CI]	Paired probability-DEF shift x 10,000 [95% CI]
GAT	42	+4.074 [+3.9, +4.2]	-0.041 [-0.060, -0.022]
GAT	43	+0.208 [+0.2, +0.3]	-0.839 [-1.892, +0.143]
GAT	44	-3.483 [-3.6, -3.4]	+0.092 [+0.068, +0.116]
GAT-Outage	42	+3.975 [+3.8, +4.1]	-0.253 [-0.323, -0.193]
GAT-Outage	43	+0.959 [+0.9, +1.0]	-0.114 [-0.343, +0.061]

Model	Seed	Attention shift x 1,000 [95% CI]	Paired probability-DEF shift x 10,000 [95% CI]
GAT-Outage	44	-3.038 [-3.1, -2.9]	+0.172 [+0.135, +0.212]

Figure 4.8 shows the paired attention and DEF shifts by checkpoint.




Figure 4.8. Checkpoints trained with seeds 42 and 43 move attention toward stale nodes and have a negative mean DEF shift. Checkpoints trained with seed 44 show the opposite directions. Error bars are 95% episode-block bootstrap intervals.

The same checkpoint pattern appears in both training regimes: four of six selected checkpoints have a positive attention shift and four have a negative mean DEF shift. Only the DEF decrease for checkpoints trained with seed 42 is separated from zero in both models; both seed-43 intervals include zero, while both checkpoints trained with seed 44 show a small increase. The effects are numerically small and vary in direction. This is not evidence for a universal claim that telemetry degradation moves attention toward stale nodes or always lowers faithfulness. The largest absolute paired probability-DEF shift is 0.0000839, equivalent to 0.00839 probability percentage points. Relative percentages are not reported because clean DEF is close to zero and would make such ratios unstable.

4.7 Analysis by action type

The combined faithfulness audit is dominated by chosen no-op actions. Across all six checkpoints, 56,550 records contain at least one valid request. Of these, 52,585 (93.0%) select no-op and 3,965 (7.0%) dispatch a request. Decisions with no valid request are excluded from these percentages. Table 4.6 reports the action composition and dispatch-only diagnostic.

Table 4.6: Action composition and dispatch-only DEF

Checkpoint	No-op (%)	Dispatch (%)	Mean dispatch-action probability	Paired DEF shift x 10,000 [95% CI]
GAT, seed 42	96.6	3.4	0.0085	0.000 [0.000, 0.000]
GAT, seed 43	88.4	11.6	0.0368	-20.358 [-32.488, -8.371]
GAT, seed 44	90.4	9.6	0.0211	-0.011 [-0.029, 0.000]
GAT-Outage, seed 42	91.4	8.6	0.0228	+0.088 [+0.047, +0.131]
GAT-Outage, seed 43	94.6	5.4	0.0127	-4.403 [-8.368, -1.285]

Checkpoint	No-op (%)	Dispatch (%)	Mean dispatch-action probability	Paired DEF shift x 10,000 [95% CI]
GAT-Outage, seed 44	92.8	7.2	0.0181	-0.044 [-0.092, +0.004]

Figure 4.9 visualizes the imbalance and the paired dispatch estimates.




Figure 4.9. (a) Absolute no-op and dispatch counts; labels give the dispatch share. (b) Mean probability assigned to the sampled action on a logarithmic scale. (c) Paired degraded-minus-clean probability DEF on the full scale. (d) Enlarged view around zero; arrows mark intervals that extend beyond the panel. Error bars are 95% episode-block bootstrap intervals.

The no-op group closely follows the combined result because it accounts for most records. The same H1 trend does not appear in the dispatch group for any checkpoint ($p = 0.44-1.00$), and the dispatch results show no consistent H4 evidence. The paired dispatch DEF shift is negative in four checkpoints, positive in one, and exactly zero in one. In four of six checkpoints its direction differs from the combined estimate. These effects are small except for the seed-43 checkpoints, and dispatch intervals are wider because fewer episodes contain scored dispatch actions. The diagnostic does not establish a separate pattern for dispatch decisions. It shows that the combined H1 and H4 findings should not be generalized to dispatch actions.

4.8 Random-trigger sensitivity analysis

The random-loss control replaces the fixed tunnel-entry trigger with an independent per-taxi trigger while retaining the same 30-second observation-layer freeze. The trigger probability is fixed across checkpoints, so realized exposure remains dependent on the trajectory induced by each policy. Random-to-tunnel exposure-rate ratios range from 0.44 to 0.98, with a median of 0.78. The comparison therefore focuses on the direction of paired effects within stale-exposed decisions rather than treating the two conditions as perfectly exposure matched. Table 4.7 reports the trigger comparison.

This sensitivity analysis is a separate 24-episode evaluation. Its repeated 30-second tunnel estimates therefore differ slightly from the 48-episode main sweep in Table 4.5; no checkpoint is retrained or reselected.

Table 4.7: Tunnel-triggered and random-loss sensitivity results

Checkpoint	Exposure tunnel/random (%)	Attention shift tunnel/random (values x 1,000)	Probability DEF shift tunnel/random (values x 10,000)
GAT, seed 42	0.580 / 0.569	+4.281 / +4.037	-0.012 / +0.004
GAT, seed 43	1.176 / 0.517	+0.231 / -0.286	-0.479 / +7.361
GAT, seed 44	0.818 / 0.532	-3.623 / -3.165	+0.020 / +0.322
GAT-Outage, seed 42	0.967 / 0.713	+4.279 / +5.453	-0.205 / -0.093
GAT-Outage, seed 43	0.584 / 0.480	+0.958 / +0.444	-0.088 / -0.014
GAT-Outage, seed 44	0.712 / 0.648	-3.063 / -4.296	+0.091 / +0.518

Figure 4.10 shows the tunnel and random-trigger sensitivity results.

Sensitivity to tunnel-triggered versus random telemetry loss




Figure 4.10. Sensitivity to tunnel and random triggers under a 30-second freeze: (a) exposure rate; (b) stale-attention shift; (c) probability-DEF shift at full scale; and (d) enlarged DEF near zero. Arrows mark off-panel intervals; error bars show 95% episode-block bootstrap intervals.

Attention-shift direction agrees between tunnel and random triggers in five of six checkpoints. Probability-DEF direction agrees in four of six. One mismatch for GAT seed 42 consists of values very close to zero. GAT seed 43 is the clear exception, but its random-trigger DEF interval is wide and crosses zero. The seed-42 positive attention pattern and seed-44 negative pattern otherwise remain visible under random triggering in both training regimes. This reduces the likelihood that the overall checkpoint dependence is produced only by the chosen tunnel location. It does not establish a location-independent effect, because the realized exposure rates and stale-node populations are not identical.

4.9 Faithfulness hypotheses

The event-aware audit uses every decision with visible stale exposure rather than a periodic sample. H1 is supported for two of three GAT checkpoints and all three GAT-Outage checkpoints. H4 is supported for all six: within-episode WAMSN-DEF correlations range from -0.199 to -0.084 for GAT and from -0.423 to -0.212 for GAT-Outage. H3 remains consistently supported. Figure 4.11 shows both the within-episode WAMSN-DEF pattern and the per-checkpoint H4 results, while Table 4.8 summarizes all hypothesis outcomes.

WAMSN-DEF relationship and H4 decision




Figure 4.11. Panel (a) divides each usable episode into four equal-count WAMSN groups and plots the corresponding mean DEF rank. Q4 is lower than Q1 for both models, although the intermediate groups are not monotonic. Lines summarize episode blocks within each model, and shaded bands show their interquartile range. Panel (b) reports mean rho and a 95% episode-block bootstrap interval. Every checkpoint supports the fourth hypothesis after Holm correction.

Table 4.8: Hypothesis outcomes across training seeds

Hypothesis	GAT seeds supporting	GAT-Outage seeds supporting	Verdict
H1: outage duration increases, DEF decreases	2/3	3/3	strong but not universal
H2: faithfulness declines faster than performance	1/3	3/3	model-dependent; exploratory
H3: outage duration increases, WAMSN increases	3/3	3/3	supported as a stale-exposure manipulation check
H4: higher WAMSN is associated with lower DEF	3/3	3/3	consistently supported within episodes
H5: GAT-Outage has weaker adverse DEF relationships than GAT	n/a	0/3	not supported in the tested configurations

For H5, every matched GAT-Outage checkpoint has more negative H1 and H4 correlations than its clean-trained GAT counterpart. The H1 comparisons for seeds 42, 43, and 44 are -0.164 versus -0.069, -0.141 versus -0.058, and -0.025 versus -0.017. The corresponding H4 comparisons are -0.423 versus -0.092, -0.411 versus -0.084, and -0.212 versus -0.199. None meets the requirement that both relationships become weaker, giving 0/3 support for H5.

This comparison does not isolate degraded training observations as the cause because GAT and GAT-Outage use different maximum training budgets and learning-rate decay horizons.

H1 and H4 are association tests within one trained policy and primarily describe no-op decisions in this action-imbalanced sample. They do not by themselves establish that AoI causes lower faithfulness or that the same relationship holds for dispatch actions. The paired clean-twin audit supplies the more direct intervention contrast, and its sign still changes with training seed. H2 is treated as exploratory because clean DEF is close to zero and the relative-rate formulation is unstable.

4.10 Result summary

The results answer the research questions as follows:

- **RQ1:** under clean telemetry, raw attention does not consistently identify decision-relevant nodes better than the type-matched random baseline;
- **RQ2:** tunnel-triggered degradation changes stale-node attention, but the direction of the paired change depends on the trained checkpoint;
- **RQ3:** longer outages consistently increase stale exposure and usually accompany lower within-policy DEF, but paired attention and DEF changes are not consistent across checkpoints, action groups, or extraction choices; and
- **RQ4:** the tested degradation-aware configuration does not produce a more consistent explanation response than the clean-trained configuration.

Seed reversals, no-op dominance, failures in the dispatch-only results, and extraction sensitivity therefore mean that raw attention cannot provide a reliable indication of data freshness or explanation faithfulness.

4.11 Demonstrative framework decision

The completed evidence was passed through the operational audit described in Section 3.10. Table 4.9 separates successful evidence collection from failed explanation-release checks.

Table 4.9: Freshness-aware explanation audit decisions

Audit check	GAT	GAT-Outage	Key evidence
Evidence integrity	PASS	PASS	3/3 checkpoint preflights passed for each model
Freshness reported separately	PASS	PASS	10,832 and 10,885 stale-exposed records reported
Action-aware controls	PASS	PASS	type matching, chosen-action protection, and LOO present
Dispatch decision relevance	FAIL	FAIL	clean/60 s CIs: GAT [-0.00367, -0.00252] / [-0.00371, -0.00252]; GAT-Outage [-0.00765, -0.00594] / [-0.00752, -0.00578]
No-op and dispatch composition	PASS	PASS	26,278/2,020 and 26,307/1,945 no-op/dispatch records
Attention-extraction stability	FAIL	FAIL	an alternative reverses the default direction in every seed
Checkpoint consistency	FAIL	FAIL	stale-attention direction is positive, positive, then negative across seeds
Deployment-action capability	N/A	N/A	sampling from the policy distribution is the audited action rule; argmax is descriptive

Audit check	GAT	GAT-Outage	Key evidence
Tunnel/random trigger robustness	INDET.	INDET.	no supported reversal, but at least one 95% interval crosses zero
Final decision	WITHHOLD	WITHHOLD	decision relevance, extraction stability, and checkpoint consistency fail

INDET. denotes INDETERMINATE: the available evidence neither passes nor fails the check.

Both models pass the checks that establish whether the evidence is complete and properly separated. Dispatch self-row margin-DEF is below the matched-random boundary under both clean and 60-second conditions, and the extraction and checkpoint checks also fail. Trigger robustness remains indeterminate because some confidence intervals cross zero. The argmax diagnostic remains an important policy limitation but is not a release gate for the sampled action rule audited here. The result is WITHHOLD for all six frozen checkpoints, so their attention maps remain internal diagnostics. This table is a demonstrative application of the post-study release rules, not a new hypothesis test or independent evidence for the main conclusion.

Chapter 5: Discussion

5.1 Answer to the central problem

The central problem is whether a graph-attention map can still be trusted when some of its vehicle data are stale. The experiment gives a two-part answer. First, the map remains available after tunnel-triggered loss freezes vehicle features at the observation boundary. Second, its displayed weights do not provide consistent evidence of either data freshness or decision relevance. The explanation can therefore outlive the data that support it, but its continued availability should not be treated as proof that it is reliable.

Three levels of evidence support this conclusion. **Exposure evidence** shows that longer outages increase the amount of attention attached to stale vehicle data. **Association evidence** shows that greater stale exposure is linked to lower DEF within all six checkpoints. **Direct paired evidence**, which compares clean and degraded observations of the same decision, is less consistent. The attention and DEF changes have the expected directions for checkpoints trained with seeds 42 and 43, but both directions reverse for seed 44 under both training regimes. The direct effects are also small. The association result therefore does not establish one general effect caused by telemetry degradation.

The analysis by action type further limits the conclusion. Chosen no-op actions account for 93.0% of scorable decisions with at least one available request, and the combined H1 and H4 patterns are not reproduced among dispatch actions. Four checkpoints also change paired-DEF direction between the combined and dispatch groups. The primary statistics describe the observed mix of sampled actions, not a universal response for both dispatch and no-op decisions.

The supporting controls make the interpretation stronger and narrower. The LOO-ranked control produces a positive DEF for every checkpoint, so the evaluator has some sensitivity to a perturbation-based ranking. However, expected top-k overlap reaches about 60% at $k=3$, which limits resolution. Taxi-only attention-LOO correlation is positive but varies substantially in strength. These checks do not show that the evaluator is perfect and attention fails. They show that, within the evaluator's measured but limited resolution, raw attention lacks reproducible evidence of explanation faithfulness. Figure 5.1 summarizes this cross-seed evidence path.

Cross-seed evidence matrix

Question	Measure	Cross-seed result	Verdict
RQ1	Clean decision relevance	Near type-matched random baseline	Not validated
RQ2	Paired stale-node attention	Positive in 4/6 trained policies	Mixed by seed
RQ3	Duration and paired DEF	Paired decline in 4/6 policies	Not consistent
RQ4	Training comparison	No consistent improvement	Not supported

Conclusion: raw attention does not provide consistent decision-relevance or degradation responses.

Figure 5.1. Cross-seed results for the four research questions. The matrix separates clean decision relevance, paired stale-node attention, duration and faithfulness evidence, and the training comparison.

5.2 What the stale-exposure result means

WAMSN rises with outage duration for every GAT and GAT-Outage checkpoint. It is useful as an operational exposure indicator: it shows when an attention map contains more weight attached to old vehicle readings.

WAMSN is not evidence that AoI causes lower faithfulness. The metric itself weights attention by normalized AoI, so part of the increase follows directly from the manipulation. H1 and H4 show that DEF tends to be lower at greater exposure within most frozen policies. The paired audit asks a stricter question: does replacing the current observation with its degraded twin change attention and DEF at the same decision? That result changes sign across checkpoints trained with different seeds. Exposure, within-policy association, and direct paired change answer different questions and should be reported separately.

5.3 Dependence on checkpoint and analysis choice

The primary aggregation gives positive stale-attention shifts for checkpoints trained with seeds 42 and 43 and negative shifts for checkpoints trained with seed 44. The values are small: approximately -0.0035 to +0.0041 of total attention mass. Paired probability-DEF shifts are also small and reproduce the same direction in both training regimes. This shared pattern across the two model conditions suggests sensitivity to the trained checkpoint, not a stable response created by outage-aware training. The between-checkpoint range is larger than the observed difference between the two training-condition means. With only three training seeds, this is a descriptive comparison of selected checkpoints. It is not a formal estimate of the sources of variation, and it does not show that the seed value itself causes the response.

The random-trigger sensitivity analysis checks whether this pattern is only an artifact of the fixed tunnel location. The same attention-shift direction appears in five of six checkpoints, and the same probability-DEF direction in four of six. The seed-42 positive attention pattern and seed-44 negative pattern appear under both triggers and both training regimes. This makes a tunnel-only explanation less likely, but it is not a complete location control. A fixed trigger

probability produces different realized exposure rates after each policy changes taxi trajectories, and the identity of the stale taxis also changes. The result supports checkpoint dependence but cannot establish one general effect of random or tunnel-triggered loss.

The paired probability-DEF effects are small and may have little operational importance. The concern is not a large universal decline. Instead, even these small changes reverse direction across independently trained checkpoints, while raw corrected DEF under clean telemetry is already near zero. Together, these results do not provide a consistent basis for trusting the explanations, but they also do not show large damage from telemetry degradation.

The aggregation audit identifies a second problem. Individual heads can reverse the sign within the same checkpoint. The six checkpoint ranges extend from -0.0109 to +0.0089 across the tested reductions. Selecting another head or layer after seeing the data could change the verbal conclusion. The declared mean remains the primary result, but the sensitivity analysis shows that it is not a unique explanation produced by the network.

Query-row choice is similarly checkpoint-dependent. The selected request row improves request-action DEF by about 0.0064 for GAT seed 43 and 0.0073 for GAT-Outage seed 43, but changes little for the other four checkpoints. One query-row change therefore cannot be assumed to correct near-zero or negative DEF. An attention map shown to an operator needs a declared and tested rule for choosing its query, layer, and heads.

The deterministic diagnostic reveals a related policy limitation. All six graph-attention checkpoints choose no-op in all 18 argmax episode evaluations, although their sampled policies outperform the declared lower bounds. This limitation does not invalidate the faithfulness calculations for sampled actions, but it means that the experiment audits decisions sampled from the policy distribution rather than a deployable deterministic dispatcher. The audit framework should evaluate the action rule intended for deployment and report any alternative action rule as a separate diagnostic.

5.4 Role of the construct-validity audit

The construct-validity audit supports the primary experiment. Request nodes are both information and actions; peer taxis are context only. Uniform random deletion therefore creates a large action-removal artifact. Type matching and chosen-request protection make the comparison fairer.

The LOO control checks a different question: can DEF respond when nodes are ranked by their own single-node margin loss? Its consistently positive result shows limited sensitivity, mainly for comprehensiveness. It is a positive control, not ground truth for the combined comprehensiveness-sufficiency score. Gradient x Input changes sign across checkpoints and is not automatically a stronger explanation in this task.

Together, these checks change the conclusion from a simple “attention is near random” statement to a more defensible one. Raw attention remains below the LOO control in all six checkpoints and has positive taxi-only agreement with LOO, but both DEF and rank agreement vary by trained policy. The aggregation and query-row audits add further sensitivity. This provides some evidence of decision relevance, but not enough to treat attention as reliable by default.

The top-k diagnostic further limits this interpretation. Attention-LOO overlap is above the expected matched-random overlap at $k=1$, but it offers no clear advantage over that baseline at $k=3$. This is consistent with reduced resolution as larger subsets overlap, not proof that the top attention node is the model's true reason or that every three-node display is random.

5.5 Why outage training did not solve the problem

The tested configuration trained with 30-second outages does not solve the explanation problem. GAT-Outage shows the same seed-dependent attention and paired DEF directions as GAT. It strengthens the H1 and H4 associations, but does not make the direct degraded-minus-clean response consistent across independently trained policies. Under the matched-seed H5 rule, both adverse correlations would need to move closer to zero. Neither does so for any of the three seeds.

These findings apply only to the tested form of robustness training. The comparison cannot show that the degraded training observations caused the difference because the two configurations have different optimization horizons. The reward also contains no term for explanation stability or faithfulness, so task-reward optimization alone is not designed to align attention with a human explanation. A future intervention would need a common training budget, independently trained models, and an explicit explanation objective.

5.6 Proposed audit framework

5.6.1 Purpose and scope

To address this risk, the thesis develops and applies a **freshness-aware explanation audit framework**. The framework leaves the trained policy unchanged and does not claim to make attention faithful. Instead, it sets rules for deciding when an attention map has enough evidence to be shown as an audited candidate explanation within a stated scope. Table 5.1 links each observed risk to the required evidence and decision rule.

Table 5.1: Freshness-aware explanation audit framework

Observed risk	Required check	Evidence	Decision rule
Attention may be attached to stale data	Report freshness separately	AoI and WAMSN	Never infer freshness from attention strength
Attention may not identify decision-relevant nodes	Use action-aware, type-matched tests	DEF, margin-DEF, and LOO	Do not claim a reason without evidence of decision relevance
Combined results may be dominated by no-op	Report no-op and dispatch groups	Action counts and action-specific DEF	Do not generalize a combined result across action types
The result may depend on attention extraction	Test the query row, layer, head, and aggregation	Sensitivity audits	Predeclare the extraction rule and report material sensitivity
One checkpoint may not represent another	Audit each checkpoint before release	Per-checkpoint results	Repeat after retraining or replacement
A pattern may not reproduce	Compare independent training runs	Results across seeds	Require a consistent conclusion across runs
The audited action may differ from deployment	Test the intended deployment action rule	Sampled-policy or argmax capability evidence	Require argmax capability only when deployment uses argmax

Observed risk	Required check	Evidence	Decision rule
A tunnel-specific pattern may not transfer	Compare tunnel and random triggers	Paired shifts and 95% confidence intervals	Withhold when supported directions conflict or remain indeterminate

The framework is intended for offline release review before an attention map is shown to a dispatcher or other operator. It should be repeated after retraining, checkpoint replacement, changes to the telemetry or graph pipeline, transfer to a new operating scenario, and during periodic model review. It is not a live freshness monitor. A deployed system must still display freshness separately and suppress or warn about explanations based on stale inputs. Figure 5.2 shows the corresponding offline audit workflow.




Figure 5.2. Saved evidence from frozen checkpoints passes through nine release checks, grouped here into six categories. The output controls whether attention may be presented as an explanation; it does not alter the model or selected action.

5.6.2 Inputs, outputs and use

The framework accepts evidence rather than raw attention alone. Table 5.2 summarizes its inputs and outputs.

Table 5.2: Audit framework inputs and outputs

Component	Contents	Role in the audit
Candidate scope	frozen checkpoints, model names, training seeds, and held-out conditions	defines exactly what the decision covers
Freshness evidence	paired clean/degraded records, AoI, stale-node masks, WAMSN, and stale-attention shift	separates data age from attention strength
Faithfulness evidence	type-matched and action-protected DEF, LOO, and dispatch-action results	tests whether highlighted nodes matter to the action
Stability evidence	action groups, extraction alternatives, checkpoint synthesis, deployment-action diagnostic, and trigger comparison	tests whether the conclusion survives relevant choices
Audit output	per-check status, per-checkpoint decision, model decision, failed reasons, and permitted use	supports a traceable explanation-release decision

The operating procedure is:

1. freeze the candidate checkpoint and evaluation protocol;
2. generate paired freshness, action-aware faithfulness, action-composition, and sensitivity evidence on held-out demand;
3. run preflight to establish that the evidence is complete and valid to analyze;
4. apply every release check to each checkpoint and then compare independent training runs;
5. present attention externally only if the result is ELIGIBLE, together with its freshness indicator and audited scope.

A preflight PASS is not an explanation-release pass. It only confirms that the evidence can be analyzed. The release output is ELIGIBLE, WITHHOLD, or INCOMPLETE. Individual checks may also be INDETERMINATE when the evidence is complete but does not support a direction. A failed or indeterminate required check produces WITHHOLD; missing evidence produces INCOMPLETE. In both cases, attention must not be presented as the reason for the action.

5.6.3 Application to this study

Applying the framework to the completed evidence gives WITHHOLD for GAT and GAT-Outage and for all six individual checkpoints. This is not a failed experiment. The audit has completed and has found that the evidence is insufficient for explanation release. In particular, the dispatch-action decision-relevance intervals do not exceed the matched-random boundary, attention extraction can reverse the stale-attention direction, the stale-attention response changes direction across training seeds, and the tunnel/random comparison is indeterminate because several confidence intervals cross zero. Every argmax diagnostic also produces zero pickups, but this is a descriptive limitation rather than a release gate for the sampled action rule.

Sampled dispatch performance provides evidence of task capability, but it cannot show that the displayed reason is current or decision-relevant.

The framework checks whether an explanation has enough evidence to be shown; it does not make the model itself more faithful. Improving the model would require an explicit explanation objective and independent evaluation, as discussed in Section 5.8.

The current study demonstrates the framework's rejection path but does not show that a real trained model can reach ELIGIBLE. The LOO positive control shows that the evaluator can reward a more informative ranking, while software tests confirm that complete passing evidence produces an ELIGIBLE decision. A future faithfulness-aware model is still needed to validate the eligibility path empirically. The current thresholds must not be weakened after seeing these results simply to obtain a passing model.

5.7 Limitations

The study uses one simulated district, 20 taxis, and 50 requests. Tunnel exposure remains scenario-dependent, although the event-aware audit scores every observed stale-exposed decision. It evaluates three independent training seeds per model. Three seeds reveal variation but cannot estimate the wider distribution of training outcomes. The selected policies pass the

training stability gates and their sampled-action evaluations outperform random and greedy lower bounds. However, every selected GAT checkpoint selects no-op throughout the small deterministic check. The study therefore cannot establish deterministic deployment capability or isolate the contribution of graph edges to sampled capability.

Clean GAT and GAT-Outage also use different maximum training budgets and learning-rate decay horizons, selected from validation-only stability diagnostics. Their matched-seed comparison describes the two fitted configurations but cannot isolate the effect of degraded training observations.

The local graph is small. Type matching makes random top-k sets overlap heavily with the explanation, which compresses DEF. LOO is based on the same single-node perturbation as comprehensiveness and is not an independent ground truth explanation. Other perturbation choices or post-hoc explainers may give different results.

The degradation layer freezes position and speed for fixed windows. Real telemetry may include delayed packets, partial failures, map-matching errors, and asynchronous recovery. The study audits attention as an explanation; it does not claim that attention is useless inside the policy computation.

The paired clean/degraded comparison supports internal validity for the tested observation-layer change because both observations come from the same simulated traffic state. Simulation and scenario-design bias remain. Road-network choice, generated demand, vehicle behavior, and tunnel placement may affect which taxis become stale and how those taxis are positioned in the local graph. In particular, a fixed tunnel can create an exposure-selection effect if its traffic conditions differ from the rest of the network. The conclusions are therefore limited to the tested SUMO setting and should not be read as population estimates for real fleets or cities.

5.8 Future work

Future experiments should add training seeds and use multiple road networks, demand levels, tunnel placements, and public mobility traces. The present random-trigger control should be extended to several calibrated exposure rates and persistent loss processes so that trigger location, exposure frequency, and outage duration can be separated. Experiments should also include disconnected-edge and message-passing ablations to establish how much the policy uses its graph channel. Delayed, intermittent, and biased telemetry should be compared with fixed freezes.

The explanation audit should also compare LOO with graph-specific post-hoc methods and evaluate objectives that directly reward explanation consistency. Any proposed improvement should be tested on held-out checkpoints, not only on more episodes from one trained model. A human-factors study could then test whether explicit freshness warnings lead operators to interpret the map more appropriately.

Chapter 6: Conclusion

This dissertation asked whether graph-attention explanations can be trusted after their supporting vehicle telemetry becomes stale. The results give a clear answer: **the attention map remains available, but this experiment does not validate its raw weights as reliable explanations under clean or degraded telemetry**. Attention strength alone does not show whether its source is current or whether the highlighted node influenced the selected action.

This is the faithfulness decoupling identified in the title. Under the tested observation-layer outages, an explanation can outlive its data because the map continues to be produced after some vehicle features have been frozen. The study does not claim that degradation always lowers a single faithfulness metric. It shows that the continued presence of an explanation is not matched by consistent evidence that the explanation is current and decision-relevant.

The four research questions are answered as follows:

- **RQ1: Under clean telemetry, do raw graph-attention weights identify decision-relevant nodes more reliably than type-matched random controls?** Not consistently. Raw-attention DEF ranges from -0.0009 to +0.0031 across the six checkpoints. LOO produces a larger positive result for all six checkpoints, while taxi-only attention-LOO rank correlation ranges from +0.088 to +0.699.
- **RQ2: Does tunnel-triggered telemetry degradation change the attention assigned to stale vehicle nodes relative to the paired clean observation?** Yes, but not in a consistent direction. Attention shifts toward stale nodes for checkpoints trained with seeds 42 and 43 and away from them for seed 44 under both training regimes. The same direction appears under random triggering in five of six checkpoints, while individual heads can still reverse it within one checkpoint.
- **RQ3: As outage duration increases, are changes in stale-node attention and decision-level faithfulness consistent across independently trained checkpoints?** No. Paired stale-node attention changes direction with the trained checkpoint. H1 is supported in five of six selected checkpoints and H4 in all six, but the combined direct paired DEF shift is negative for four checkpoints and positive for two. The effect is small and reverses for the checkpoints trained with seed 44. Moreover, 93.0% of scorable decisions with at least one available request are no-op, and the H1/H4 patterns are not reproduced in the dispatch group.
- **RQ4: In the tested configurations, does degradation-aware training produce more consistent attention and faithfulness responses under telemetry degradation than clean training?** No. GAT-Outage shows the same seed-dependent paired direction as GAT and provides no consistent improvement. Because the configurations use different training budgets, this comparison does not isolate degraded training observations as the cause.

This interpretation depends on the construct-validity audit. Deleting a request can also delete an action, while deleting a peer taxi only hides context. Type-matched and action-protected controls remove much of this artifact. The positive LOO result shows that DEF has some perturbation sensitivity, but the high expected top-k overlap limits its resolution. The conclusion draws on several checks rather than one near-zero number.

The 30-second random-trigger sensitivity analysis reduces concern that the checkpoint-dependent pattern is produced only by the selected tunnel. It does not remove the need for multiple tunnel placements or real telemetry traces, because realized exposure and the stale-node population still differ between trigger mechanisms.

The performance results also have an important limit. Sampled policies outperform the basic lower bounds, but all six graph-attention checkpoints choose no-op in 18 of 18 deterministic diagnostic episodes. The study audits explanations for actions drawn from the policy distribution; it does not demonstrate a deployable argmax dispatcher.

Taken together, the research questions separate three findings. Stale-data exposure rises with outage duration. Higher exposure is associated with lower DEF within the tested checkpoints. However, the direct paired effect of degradation is small and changes direction across checkpoints. These results do not contradict one another: exposure, association, and direct paired change answer different questions. They show why freshness and faithfulness must be reported separately.

The freshness-aware explanation audit framework addresses this gap. It reports freshness and action composition separately, uses action-aware and type-matched tests, checks the attention extraction rule, audits every candidate checkpoint, and requires consistency across independent training runs. Stable dispatch output is not an explanation test.

Its demonstrative application gives WITHHOLD for GAT, GAT-Outage, and all six frozen checkpoints. This means that the audit completed but the attention maps did not meet the release conditions. An ELIGIBLE model is not required to support the present finding: correctly withholding unsupported explanations is the purpose of the framework. However, the study validates only the rejection path on trained models. A future model with an explicit explanation objective is needed to test the eligibility path on held-out evidence.

An ELIGIBLE result would permit attention to be shown only as an audited candidate explanation within the stated scope, with freshness displayed separately. It would not prove a complete causal explanation.

The framework is designed to reduce operational risk, not to remove faithfulness decoupling from the model. Attention remains an internal diagnostic unless the full audit provides consistent evidence.

All six objectives were completed: the paired benchmark and audit pipeline were implemented, the four research questions were evaluated, and the collected evidence produced a WITHHOLD decision for every tested checkpoint. Completion refers to execution of the planned study, not support for every hypothesis or achievement of an ELIGIBLE result.

In summary, under telemetry degradation in graph-attention MARL fleet dispatch, an explanation can remain visible after part of its supporting observation has become stale. Because the displayed attention does not provide consistent evidence of freshness or decision relevance, it should not be released as an explanation without a separate audit. The contribution of this study is an evidence-based way to make that release decision, not a claim that raw attention is inherently trustworthy or that degradation always reduces DEF.

List of Publications

No publication is claimed in this dissertation.

References

- Abnar, S. and Zuidema, W. (2020) 'Quantifying attention flow in transformers', in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4190-4197. Available at: <https://doi.org/10.18653/v1/2020.acl-main.385>.
- Adebayo, J. et al. (2018) 'Sanity checks for saliency maps', *Advances in Neural Information Processing Systems*, vol. 31. Available at: <https://proceedings.neurips.cc/paper/2018/hash/294a8ed24b1ad22ec2e7efea049b8737-Abstract.html>.
- Alvarez-Melis, D. and Jaakkola, T.S. (2018) 'On the robustness of interpretability methods', *arXiv preprint arXiv:1806.08049*. Available at: <https://doi.org/10.48550/arXiv.1806.08049>.
- Amara, K. et al. (2022) 'GraphFramEx: Towards systematic evaluation of explainability methods for graph neural networks', in *Proceedings of the First Learning on Graphs Conference, Proceedings of Machine Learning Research*, vol. 198, pp. 44:1-44:23. Available at: <https://proceedings.mlr.press/v198/amara22a.html>.
- Bekkemoen, Y. (2024) 'Explainable reinforcement learning (XRL): A systematic literature review and taxonomy', *Machine Learning*, vol. 113, no. 1, pp. 355-441. Available at: <https://doi.org/10.1007/s10994-023-06479-7>.
- Bouteiller, Y., Ramstedt, S., Beltrame, G., Pal, C.J. and Binas, J. (2021) 'Reinforcement learning with random delays', in *Proceedings of the 9th International Conference on Learning Representations*. Available at: <https://openreview.net/forum?id=QFYnKlBJYR>.
- DeYoung, J. et al. (2020) 'ERASER: A benchmark to evaluate rationalized NLP models', in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 4443-4458. Available at: <https://doi.org/10.18653/v1/2020.acl-main.408>.
- Greydanus, S. et al. (2018) 'Visualizing and understanding Atari agents', in *Proceedings of the 35th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 80, pp. 1792-1801. Available at: <https://proceedings.mlr.press/v80/greydanus18a.html>.
- Heuillet, A., Couthouis, F. and Diaz-Rodriguez, N. (2021) 'Explainability in deep reinforcement learning', *Knowledge-Based Systems*, vol. 214, article 106685. Available at: <https://doi.org/10.1016/j.knosys.2020.106685>.
- Hooker, S. et al. (2019) 'A benchmark for interpretability methods in deep neural networks', *Advances in Neural Information Processing Systems*, vol. 32. Available at: <https://proceedings.neurips.cc/paper/2019/hash/fe4b8556000d0f0cae99daa5c5c5a410-Abstract.html>.
- Hu, Y., Feng, S. and Li, S. (2025) 'BMG-Q: Localized bipartite match graph attention Q-learning for ride-pooling order dispatch', *IEEE Transactions on Intelligent Transportation Systems*, vol. 26, no. 10, pp. 15407-15421. Available at: <https://doi.org/10.1109/TITS.2025.3595653>.
- Iqbal, S. and Sha, F. (2019) 'Actor-attention-critic for multi-agent reinforcement learning', in *Proceedings of the 36th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 97, pp. 2961-2970. Available at: <https://proceedings.mlr.press/v97/iqbal19a.html>.
- Jacovi, A. and Goldberg, Y. (2020) 'Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness?', in *Proceedings of the 58th Annual*

- Meeting of the Association for Computational Linguistics*, pp. 4198-4205. Available at: <https://doi.org/10.18653/v1/2020.acl-main.386>.
- Jain, S. and Wallace, B.C. (2019) 'Attention is not explanation', in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, vol. 1, pp. 3543-3556. Available at: <https://aclanthology.org/N19-1357/>.
- Kaul, S., Yates, R. and Gruteser, M. (2012) 'Real-time status: How often should one update?', in *Proceedings of IEEE INFOCOM*, pp. 2731-2735. Available at: <https://doi.org/10.1109/INFOCOM.2012.6195689>.
- Lin, K. et al. (2018) 'Efficient large-scale fleet management via multi-agent deep reinforcement learning', in *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1774-1783. Available at: <https://doi.org/10.1145/3219819.3219993>.
- Liotet, P., Maran, D., Bisi, L. and Restelli, M. (2022) 'Delayed reinforcement learning by imitation', in *Proceedings of the 39th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 162, pp. 13528-13556. Available at: <https://proceedings.mlr.press/v162/liotet22a.html>.
- Lipton, Z.C. (2018) 'The mythos of model interpretability', *Queue*, vol. 16, no. 3, pp. 31-57. Available at: <https://doi.org/10.1145/3236386.3241340>.
- Liu, Y. et al. (2022) 'Rethinking attention-model explainability through faithfulness violation test', in *Proceedings of the 39th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 162, pp. 13807-13824. Available at: <https://proceedings.mlr.press/v162/liu22i.html>.
- Lopez, P.A. et al. (2018) 'Microscopic traffic simulation using SUMO', in *Proceedings of the 21st IEEE International Conference on Intelligent Transportation Systems*, pp. 2575-2582. Available at: <https://doi.org/10.1109/ITSC.2018.8569938>.
- Lowe, R. et al. (2017) 'Multi-agent actor-critic for mixed cooperative-competitive environments', *Advances in Neural Information Processing Systems*, vol. 30. Available at: <https://proceedings.neurips.cc/paper/2017/hash/68a9750337a418a86fe06c1991a1d64c-Abstract.html>.
- Luo, D. et al. (2020) 'Parameterized explainer for graph neural network', *Advances in Neural Information Processing Systems*, vol. 33, pp. 19620-19631. Available at: <https://proceedings.neurips.cc/paper/2020/hash/e37b08dd3015330dcbb5d6663667b8b8-Abstract.html>.
- Madumal, P. et al. (2020) 'Explainable reinforcement learning through a causal lens', *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 34, no. 3, pp. 2493-2500. Available at: <https://doi.org/10.1609/aaai.v34i03.5631>.
- Milani, S. et al. (2024) 'Explainable reinforcement learning: A survey and comparative review', *ACM Computing Surveys*, vol. 56, no. 7, pp. 1-36. Available at: <https://doi.org/10.1145/3616864>.
- Mott, A. et al. (2019) 'Towards interpretable reinforcement learning using attention augmented agents', *Advances in Neural Information Processing Systems*, vol. 32. Available at: <https://proceedings.neurips.cc/paper/2019/hash/e9510081ac30ffa83f10b68cde1cac07-Abstract.html>.
- Puiutta, E. and Veith, E.M.S.P. (2020) 'Explainable reinforcement learning: A survey', in Holzinger, A. et al. (eds.) *Machine Learning and Knowledge Extraction*. Lecture Notes in Computer Science, vol. 12279. Cham: Springer, pp. 77-95. Available at: https://doi.org/10.1007/978-3-030-57321-8_5.

- Qin, Z., Zhu, H. and Ye, J. (2022) 'Reinforcement learning for ridesharing: An extended survey', *Transportation Research Part C: Emerging Technologies*, vol. 144, article 103852. Available at: <https://doi.org/10.1016/j.trc.2022.103852>.
- Rashid, T. et al. (2020) 'Monotonic value function factorisation for deep multi-agent reinforcement learning', *Journal of Machine Learning Research*, vol. 21, no. 178, pp. 1-51. Available at: <https://jmlr.org/papers/v21/20-081.html>.
- Ribeiro, M.T., Singh, S. and Guestrin, C. (2016) 'Why should I trust you? Explaining the predictions of any classifier', in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135-1144. Available at: <https://doi.org/10.1145/2939672.2939778>.
- Scarselli, F. et al. (2009) 'The graph neural network model', *IEEE Transactions on Neural Networks*, vol. 20, no. 1, pp. 61-80. Available at: <https://doi.org/10.1109/TNN.2008.2005605>.
- Serrano, S. and Smith, N.A. (2019) 'Is attention interpretable?', in *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pp. 2931-2951. Available at: <https://doi.org/10.18653/v1/P19-1282>.
- Sha, J. et al. (2026) 'A multi-agent reinforcement learning scheduling algorithm integrating state graph and task graph structural modeling for ride-sharing dispatching', *Scientific Reports*, vol. 16, article 5461. Available at: <https://doi.org/10.1038/s41598-026-35004-8>.
- Tabassi, E. (2023) *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. Gaithersburg, MD: National Institute of Standards and Technology. Available at: <https://doi.org/10.6028/NIST.AI.100-1>.
- Vaswani, A. et al. (2017) 'Attention is all you need', *Advances in Neural Information Processing Systems*, vol. 30, pp. 5998-6008. Available at: <https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html>.
- Velickovic, P. et al. (2018) 'Graph attention networks', in *Proceedings of the International Conference on Learning Representations*. Available at: <https://openreview.net/forum?id=rJXMpikCZ>.
- Wang, J. et al. (2025) 'CoopRide: Cooperate all grids in city-scale ride-hailing dispatching with multi-agent reinforcement learning', in *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, vol. 1, pp. 1457-1468. Available at: <https://doi.org/10.1145/3690624.3709205>.
- Wiegrefe, S. and Pinter, Y. (2019) 'Attention is not not explanation', in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pp. 11-20. Available at: <https://doi.org/10.18653/v1/D19-1002>.
- Yates, R.D. et al. (2021) 'Age of information: An introduction and survey', *IEEE Journal on Selected Areas in Communications*, vol. 39, no. 5, pp. 1183-1210. Available at: <https://doi.org/10.1109/JSAC.2021.3065072>.
- Ying, R. et al. (2019) 'GNNExplainer: Generating explanations for graph neural networks', *Advances in Neural Information Processing Systems*, vol. 32, pp. 9240-9251. Available at: <https://proceedings.neurips.cc/paper/2019/hash/d80b7040b773199015de6d3b4293c8f-f-Abstract.html>.
- Yu, C. et al. (2022) 'The surprising effectiveness of PPO in cooperative multi-agent games', *Advances in Neural Information Processing Systems*, vol. 35, pp. 24611-24624. Available at: <https://doi.org/10.52202/068431-1787>.

- Yuan, H. et al. (2023) 'Explainability in graph neural networks: A taxonomic survey', *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 5, pp. 5782-5799. Available at: <https://doi.org/10.1109/TPAMI.2022.3204236>.
- Yuan, H. et al. (2021) 'On explainability of graph neural networks via subgraph explorations', in *Proceedings of the 38th International Conference on Machine Learning, Proceedings of Machine Learning Research*, vol. 139, pp. 12241-12252. Available at: <https://proceedings.mlr.press/v139/yuan21c.html>.

Appendices

Appendix A: Experiment Reproduction

The complete reproduction procedure is provided in docs/REPRODUCE_EXPERIMENTS.md. Compact citable outputs are stored in the version-controlled results directory identified in that guide.

Appendix B: Supporting Artifacts

The repository contains the experiment configuration, source code, validation selections, preflight reports, per-seed analyses, summary tables and figure sources used in this dissertation. Raw cells and checkpoints remain outside the document because of their size.